

Networks and Systems Security



Dr Sudipta Saha

<https://www.iitbbs.ac.in/profile.php/sudipta/>

Decentralized and Smart Systems
Research Group (DSSRG)


<https://sites.google.com/iitbbs.ac.in/dssrg>

Computer Science & Engineering, School of Electrical Sciences
Indian Institute of Technology Bhubaneswar

Intro to *Collision Detection*


A transmitting node *can* receive its *own* transmission:

Node can receive own transmission

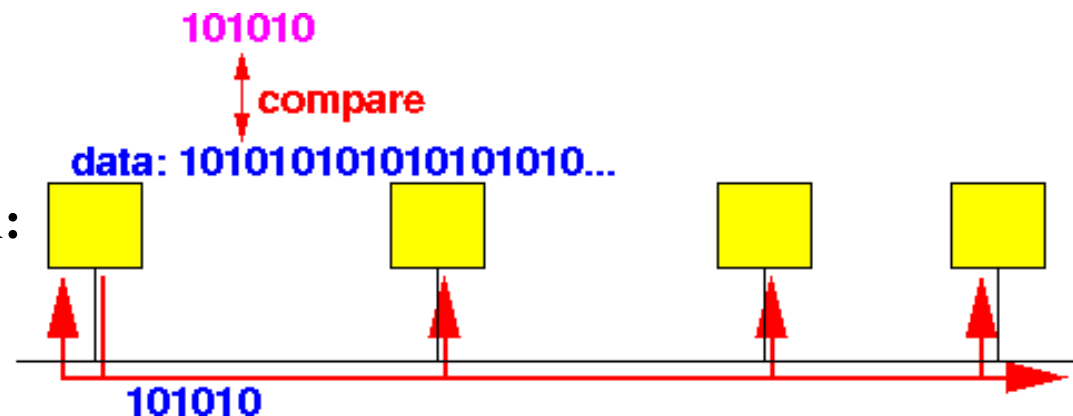


How to detect a collision:

- While transmitting, the sender receives its *own* transmission:



- *Compare* the data received against the data transmitted:



Collision will
be detected through:
Bit errors




Collisions in the CSMA protocols

- ***Every* variant** of the **CSMA protocols** can have **collisions**




Collision in Non-Persistent CSMA

- In *non-persistent CSMA*, a **collision** happens when 2 nodes pick the *same* random number:



Collision in 1-Persistent CSMA


- In *1*-persistent CSMA, a **collision** happens when 2 nodes starts sensing during the *same* transmission:



Result of collisions:


- The **transmission (frames)** involved in a **collision** are *lost*:

-



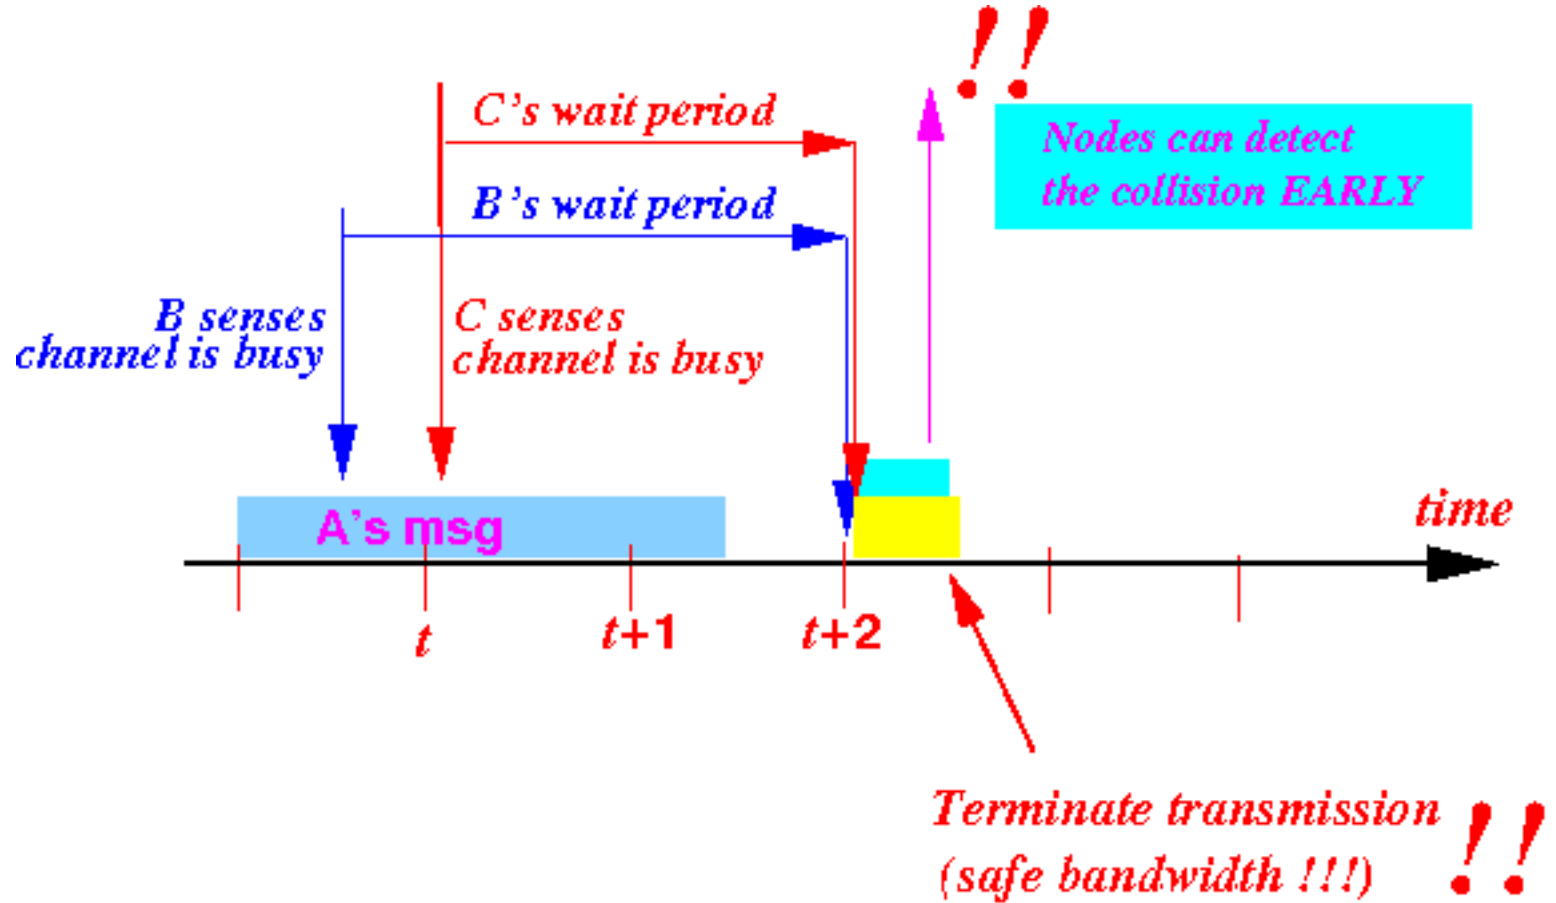
How CD helps (Collision Detection)

- *Improving* the CSMA protocol with *Collision Detection*
- The nodes can detect a collision very *early*:
-



How CD helps

- The nodes can *terminate* the transmission prematurely:




-
- **Result:**
 - ***Less* transmission time are *wasted* in a *useless* collision**
 -




Why CSMA/CD uses unslotted transmissions

- The *collision* period is *very* short compared to a slot time:



Why CSMA/CD uses unslotted transmissions

- *Slotted* transmissions will *waste* the *remainder* of the slot:




- *Unslotted* transmissions to *minimize* wastage of transmission time



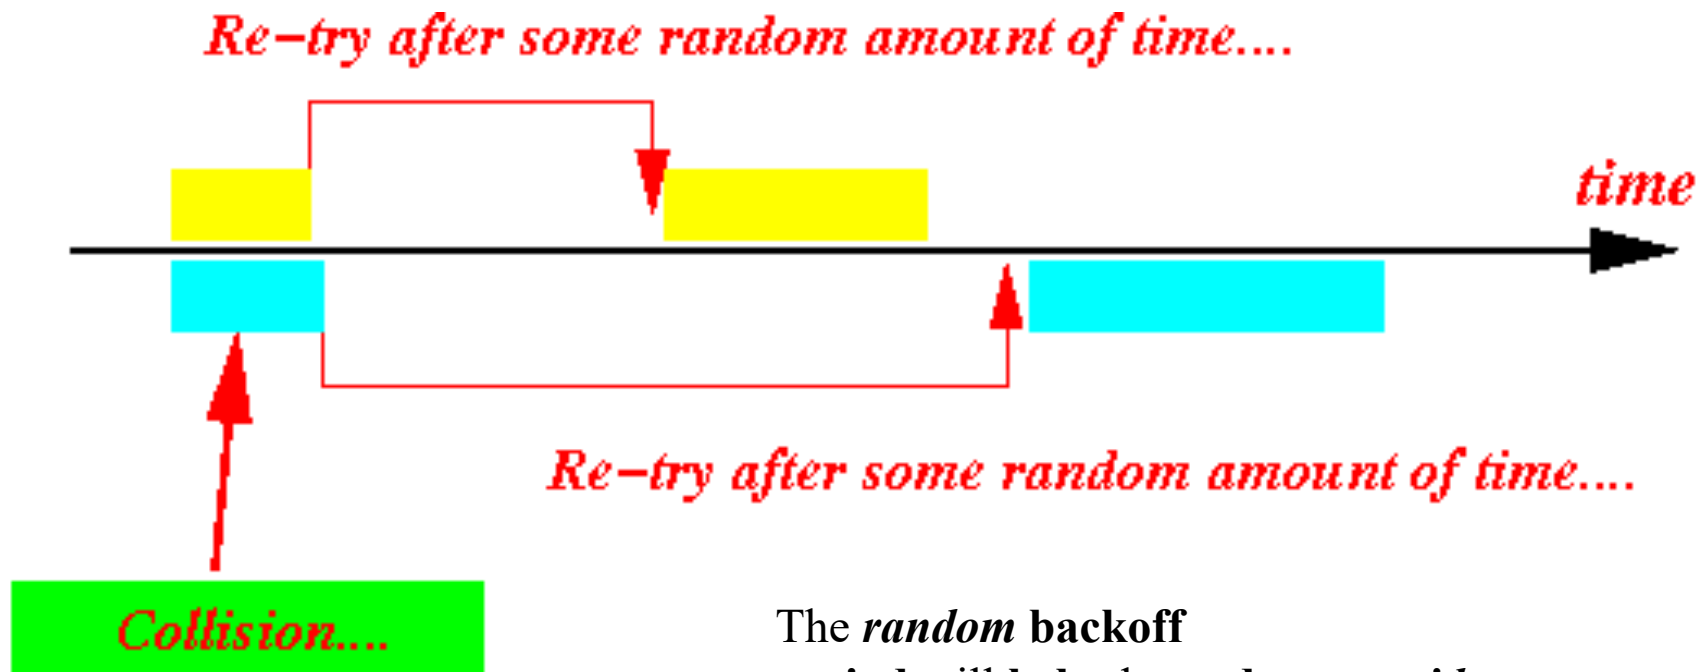
What happens *after* a node detects a *collision* ?

- *Action taken after a positive collision detection --- the CSMA/CD protocol*
- Colliding frames get lost ...



What to do when a node detects a collision:

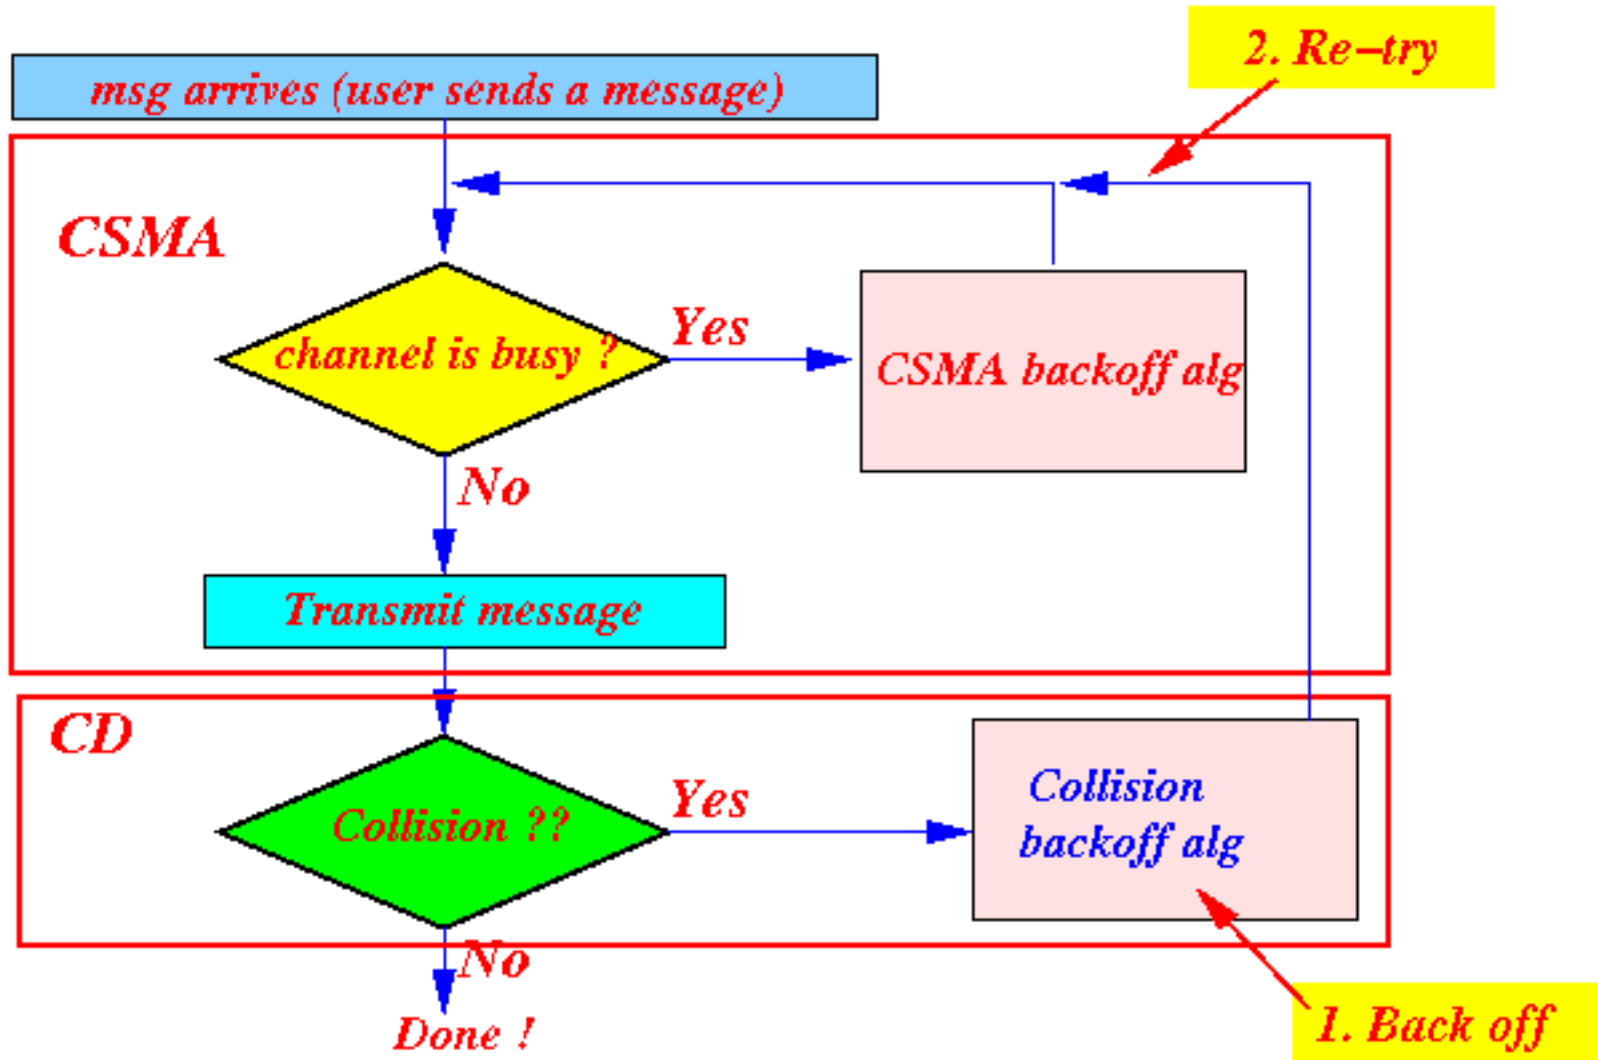
- *Back off* (= wait) for a random period
- Re-try again



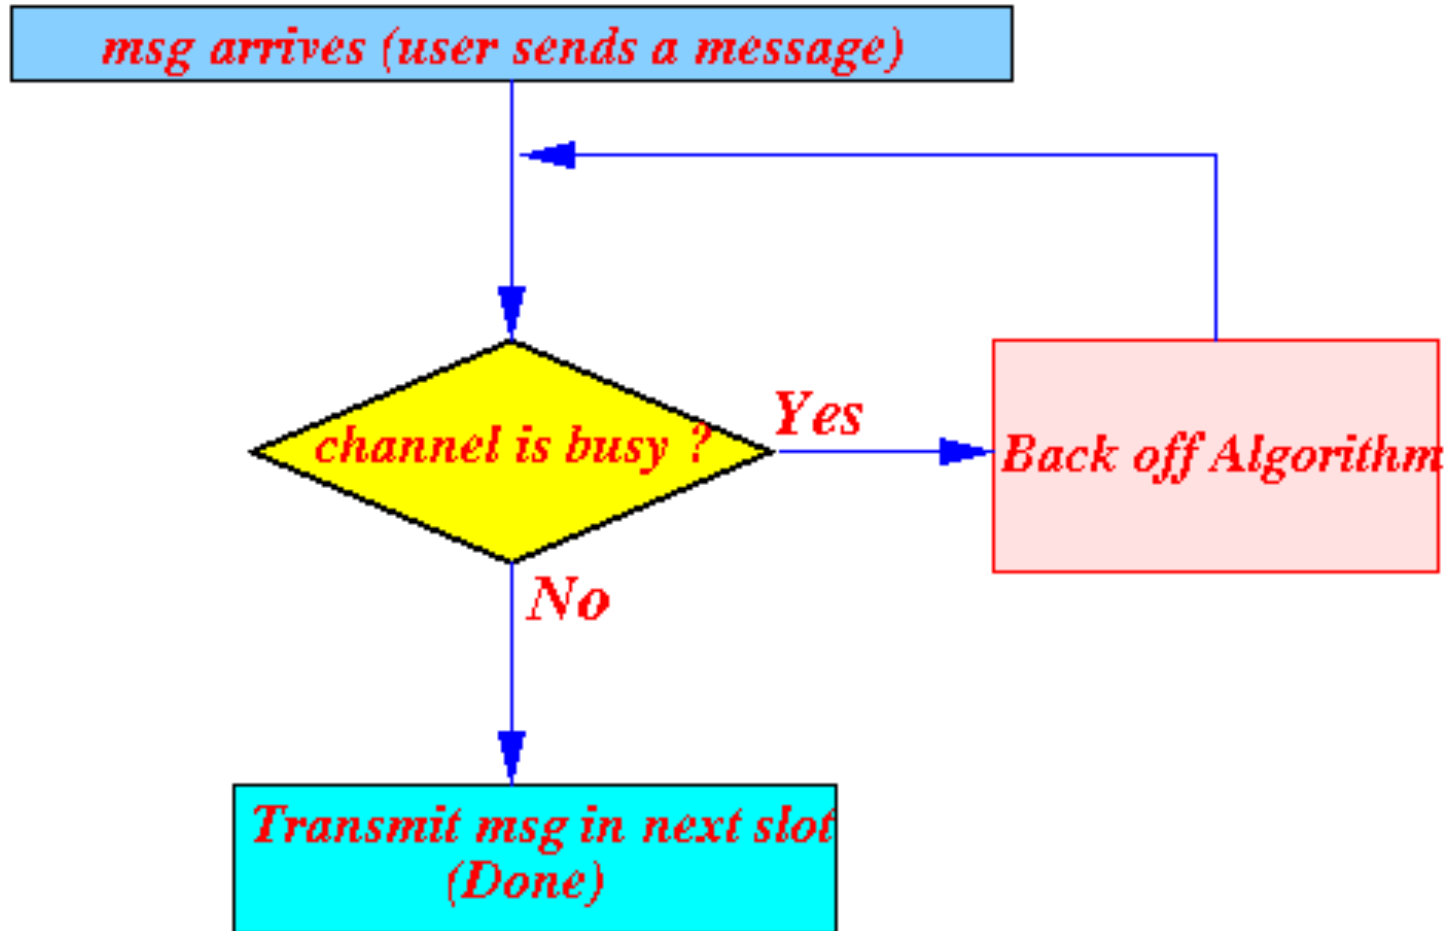
The *random* backoff period will help the nodes to *avoid* a new collision when they attempt to re-transmit their frame



The CSMA protocol with collision detection (CD)



Compare that to the flow chart for CSMA:



Two issues to address now

1. What is the unit of the back-off time?

[The answer you know already]


2. How long a node needs to sense to detect a collision has happened or not?

• We discuss these two issues next ...



The time unit used in the Back Off Algorithm in CSMA/CD

- **Recall: End-to-end propagation delay**
- **End-to-end propagation delay = amount of time that it takes for the electric signals to travel from one-end of the network to the *other* end of the network:**
-




Notation: τ = end-to-end propagation delay



- Property of the end-to-end propagation delay (τ)
- *All* nodes on the LAN will hear a node's transmission *after* τ sec after the start of the transmission
- A (any) node on the LAN *starts* transmitting at time t :
-

Current time:


t



- At time $(t + \tau)$ sec, *every* node on the LAN will hear (= detect) the transmission:

- *Current time:*

$t + \tau$



Back off time period:

- **Back off time duration** $= r \times \tau$
- $(r = \text{a random number})$
- We choose τ to be the "time unit" of the back off algorithm
- We see the reason why next...
-

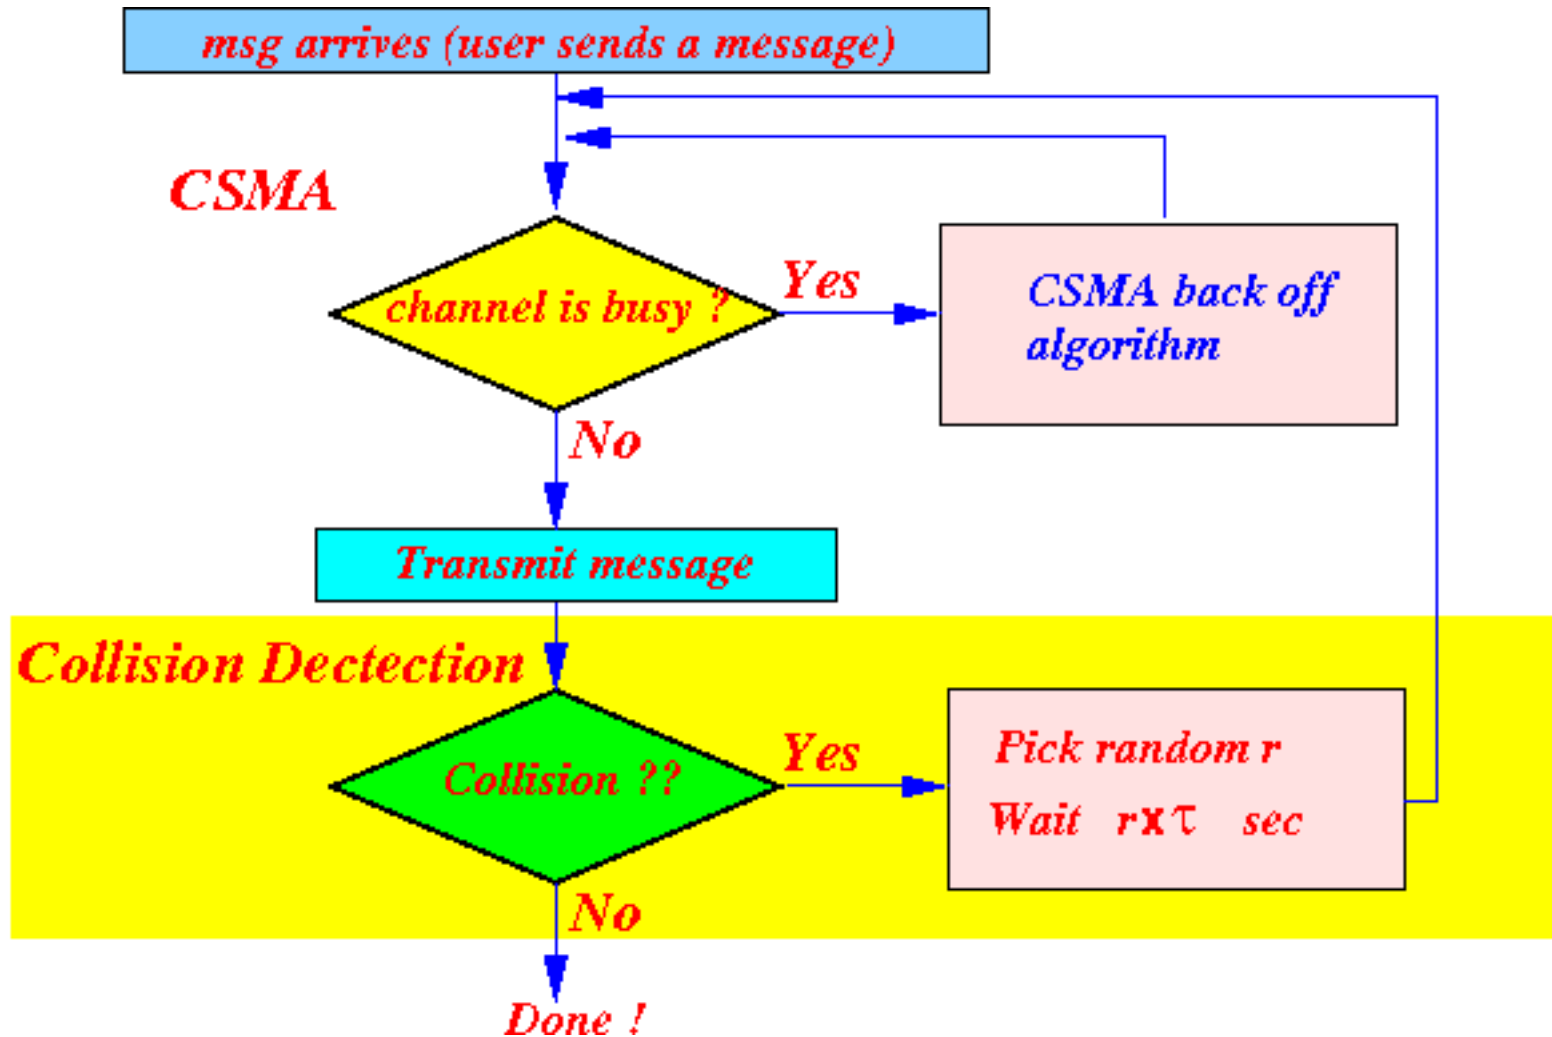


The generic Back Off Algorithm in CSMA/CD

- Pick a random integer r
- Wait $r \times \tau$ seconds
- (Start the transmission (CSMA/CD) protocol from the start)
- If 2 different nodes pick *different* random integers:
 - The **back off periods** used by these **nodes** will **ensure** (**guarantee**) the *later* starting node will **hear** (= **detect** !) the **transmission** of the *earlier* starting node when the *later* starting node **re-tries** its **transmission** !!!
- This will **allow** the *later* starting node **avoid** **colliding** with the *earlier* starting node !!!



Updated (improved) flow chart for CSMA/CD:



How long a node needs to sense for collision?

- A node does not need to perform collision detection forever
- A node can stop sensing for collision after $2 \times \tau$ time

Claim


- ***If*** node has transmitted for $2 \times \tau$ (τ = end-to-end-delay), and it **did *not*** detect a collision,
- ***then***
- The **transmission** will **not** suffer a **collision** for ***certain*** !!




Proof:

- In **CSMA/CD**, a **node** will **not** transmit if it hears an *on-going* transmission:

-




- We consider the *worst case* scenario where:
- a **collision** takes place at the *last* possible moment
- The node *A* at one end of the network start transmitting at so *arbitrary* time *t*:
-



- *Just* before the *last* node (C) hears A's transmission, it starts to transmit:


-



This will happen just *before* $t + \tau$ sec

There is now a collision !!

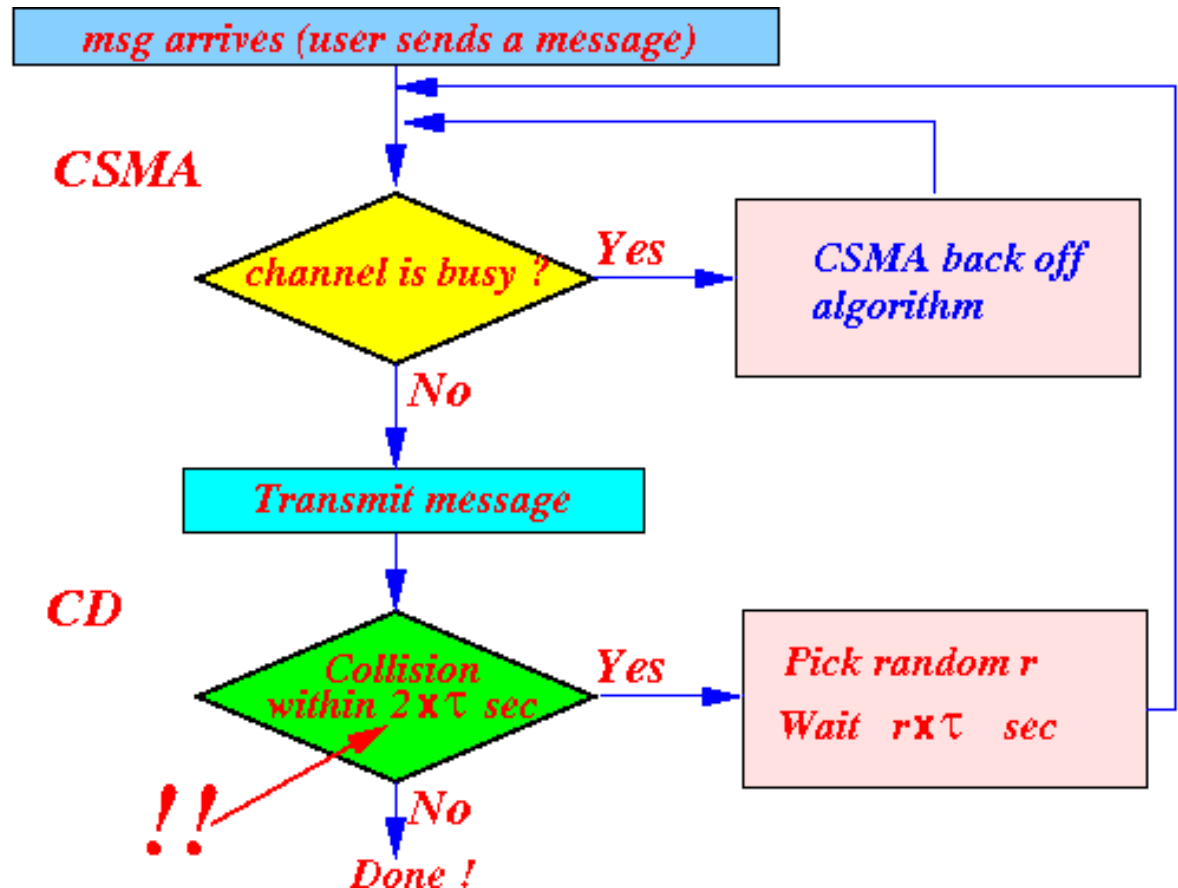
- The node *A* will **detect** the *collision* when *C*'s transmission reaches *A*:



- This will **happen** at time $t + 2 \times \tau$ sec


Therefore:

*If the node A has **not** detected a collision after $2 \times \tau$ time, there **will not** be any **future** collision !!!*




Jamming to ensure collision detection

- What is the duration of a collision,
- Suppose node *A* starts a transmission:




- Node *B* starts to **transmit** at Δ time *before* *A*'s **transmission** arrives at *B*:




- Suppose node *B* stops transmitting when *B* hears *A*'s transmission:

-



- Then:


-




the *duration* of the **collision** is equal to: Δ sec

The duration of collision can be *very* short

- If a node (*B*) starts a transmission *just* before a transmission *arrives*:



- the *duration* of collision will be *very* short



Collision detection may fail...

- ***Short* collision periods generates *few* bit errors and *may* escape detection !!!**
- **How to improve collision detection?**
-



Making sure that a **collision** is *easily* detectable:


- When a *transmitting* node has detected a **collision**, it will:
 1. **Transmit a (short) jam signal**
 2. **Stop transmitting**
- The **jam signal** is used to **help** the *other* transmitting nodes detect the **collision**



- **Example:**


- **Node *B* starts just before *A*'s transmission arrives:**

-




- **Node *B* has detected the collision:**

-



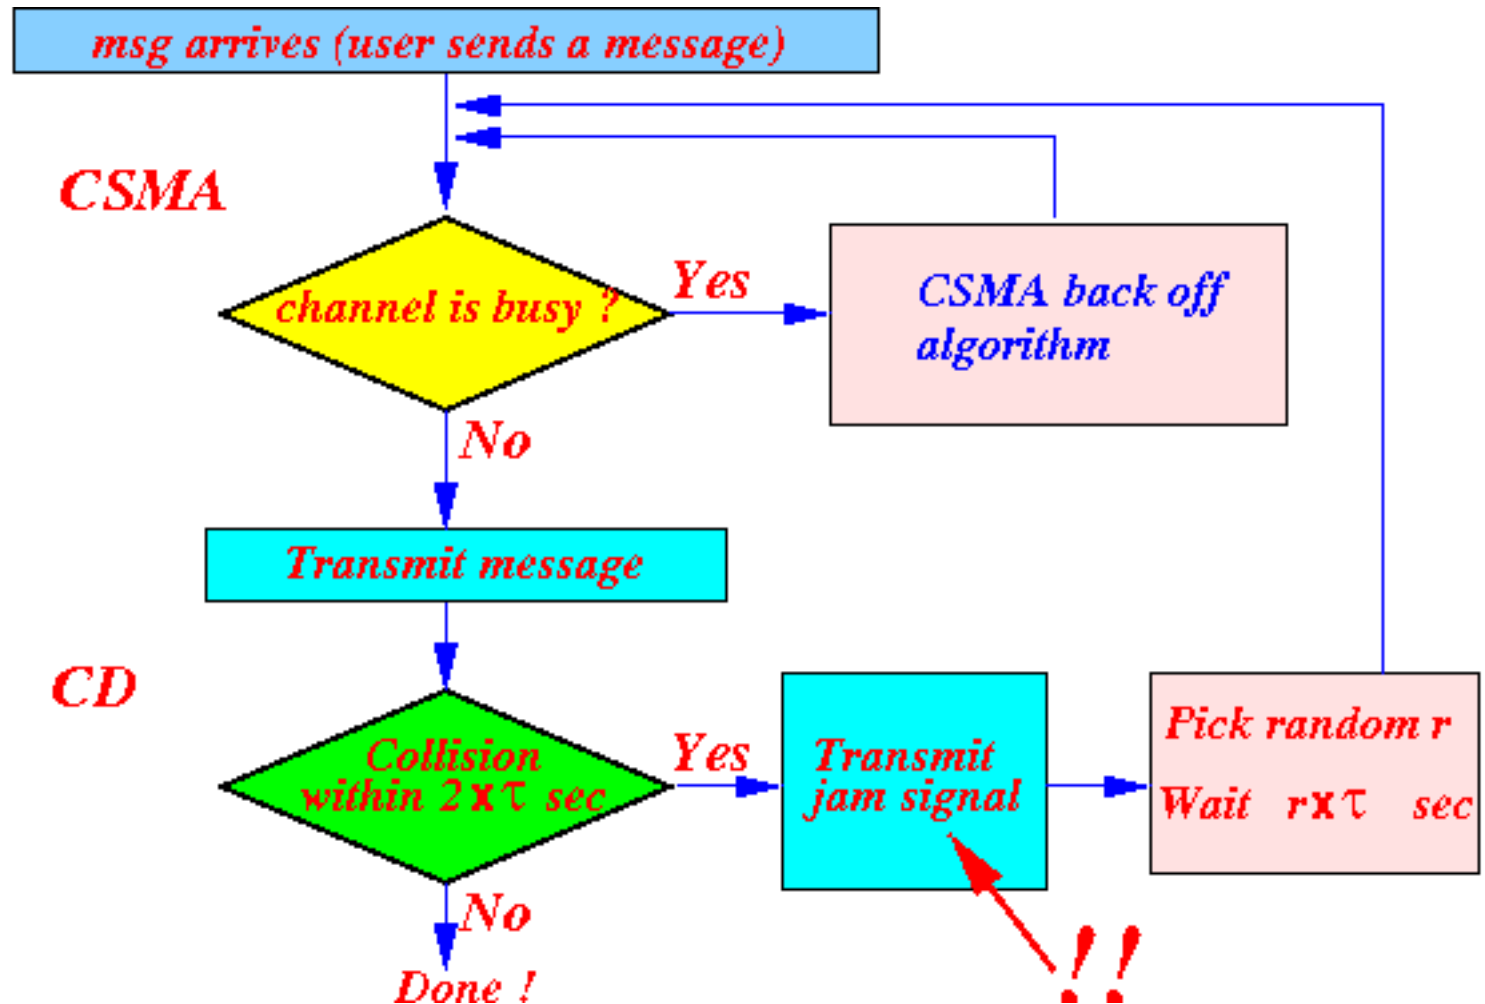
Very short collision period

- **Node *B*** transmits a (short) **jam signal** to improve **collision detection**:



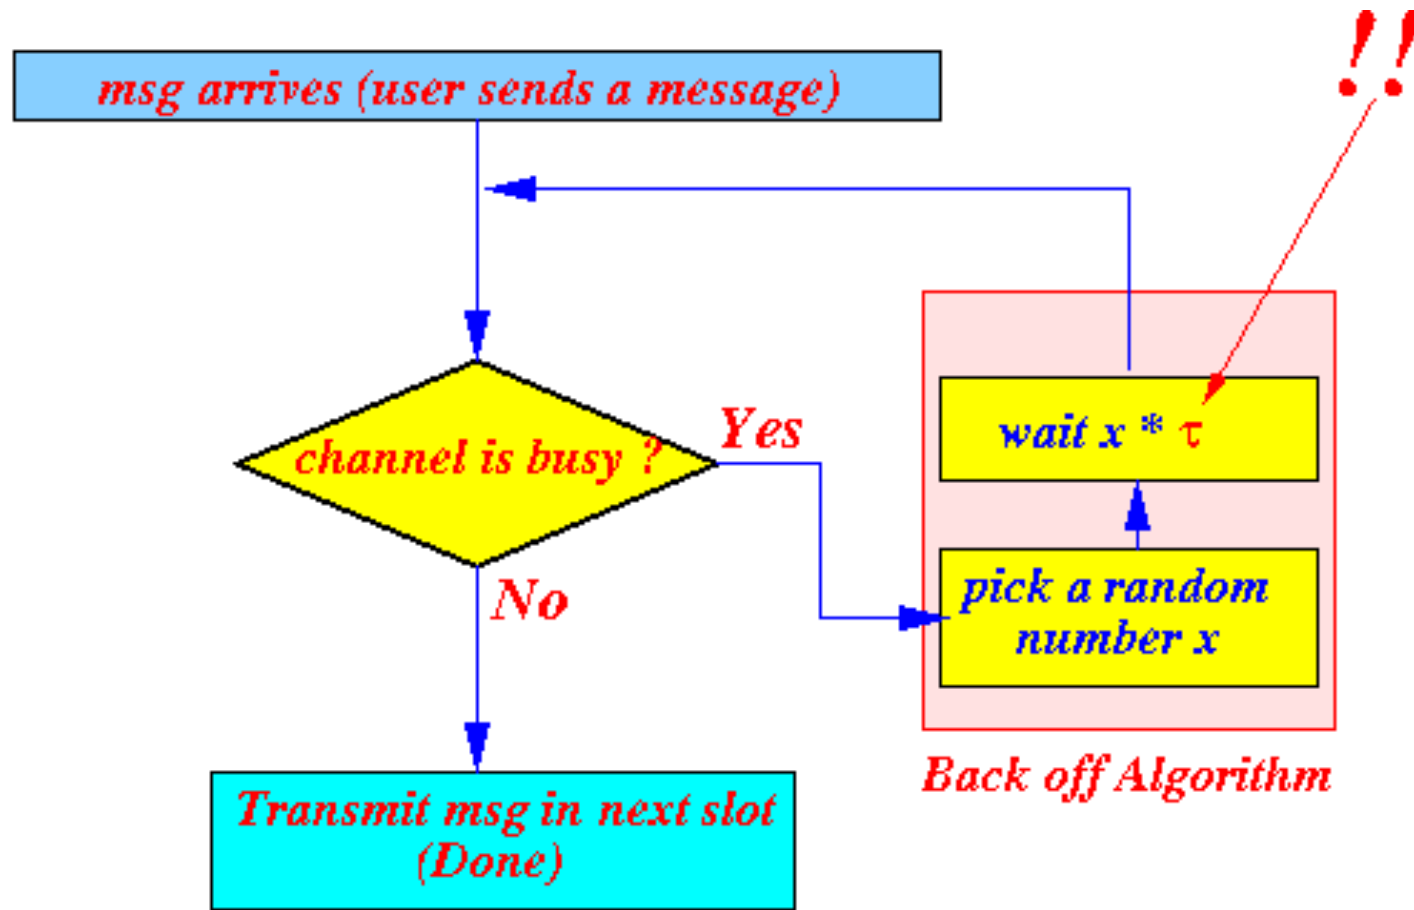
*Longer collision period !!!
(Easier to detect !!!)*

The final CSMA/CD protocol (chart)



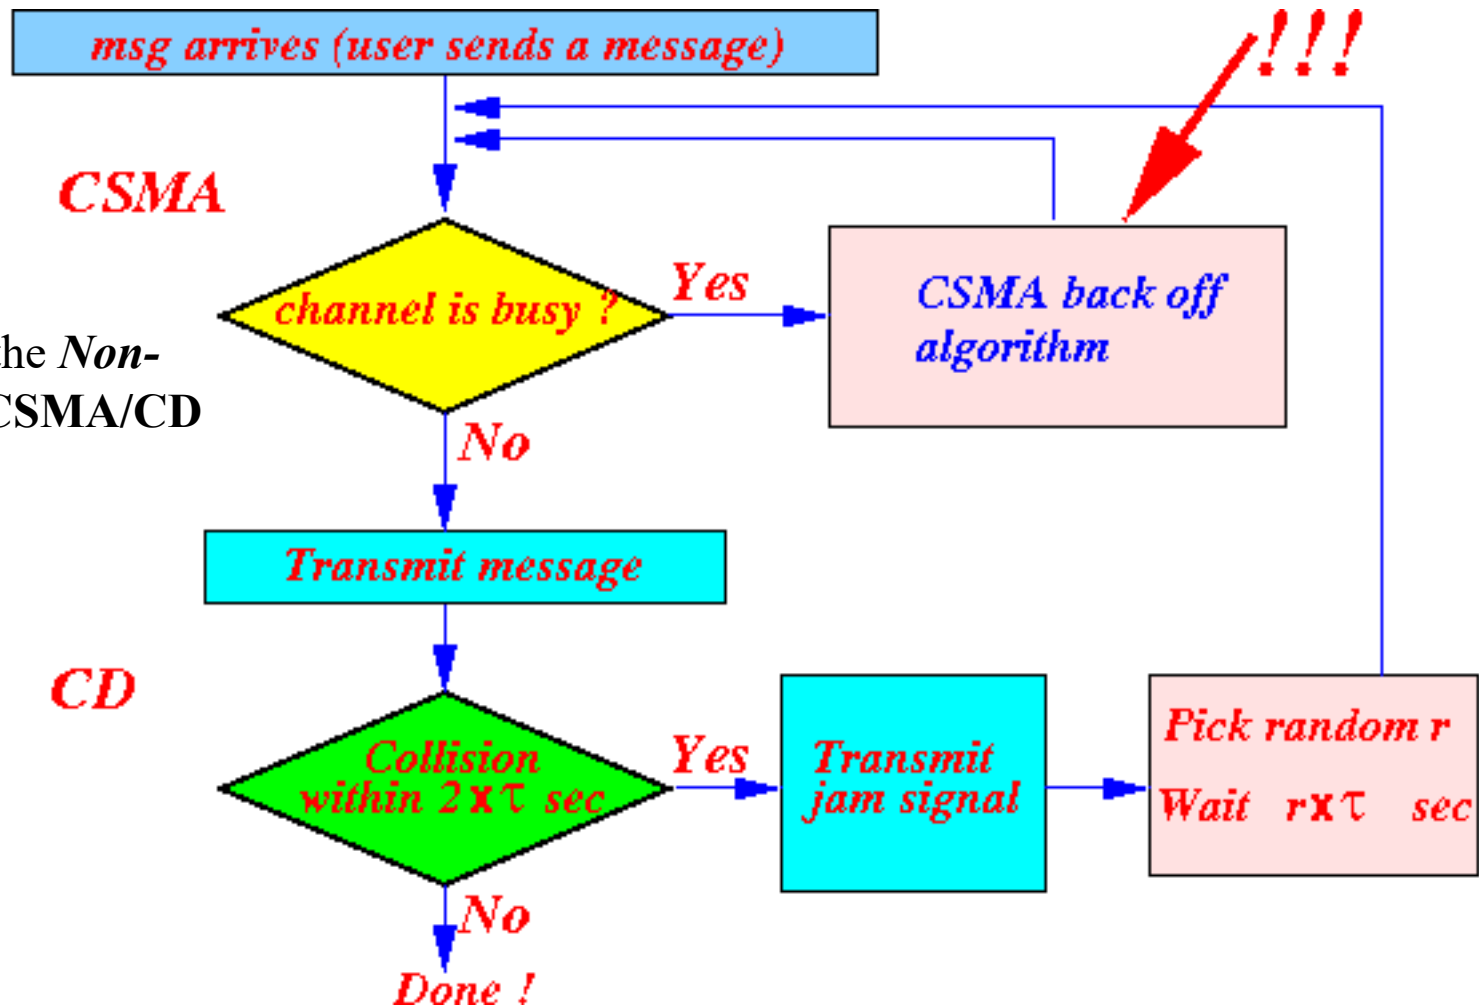
Variants of the CSMA/CD protocols

- Recall the *non-persistent* CSMA protocol:



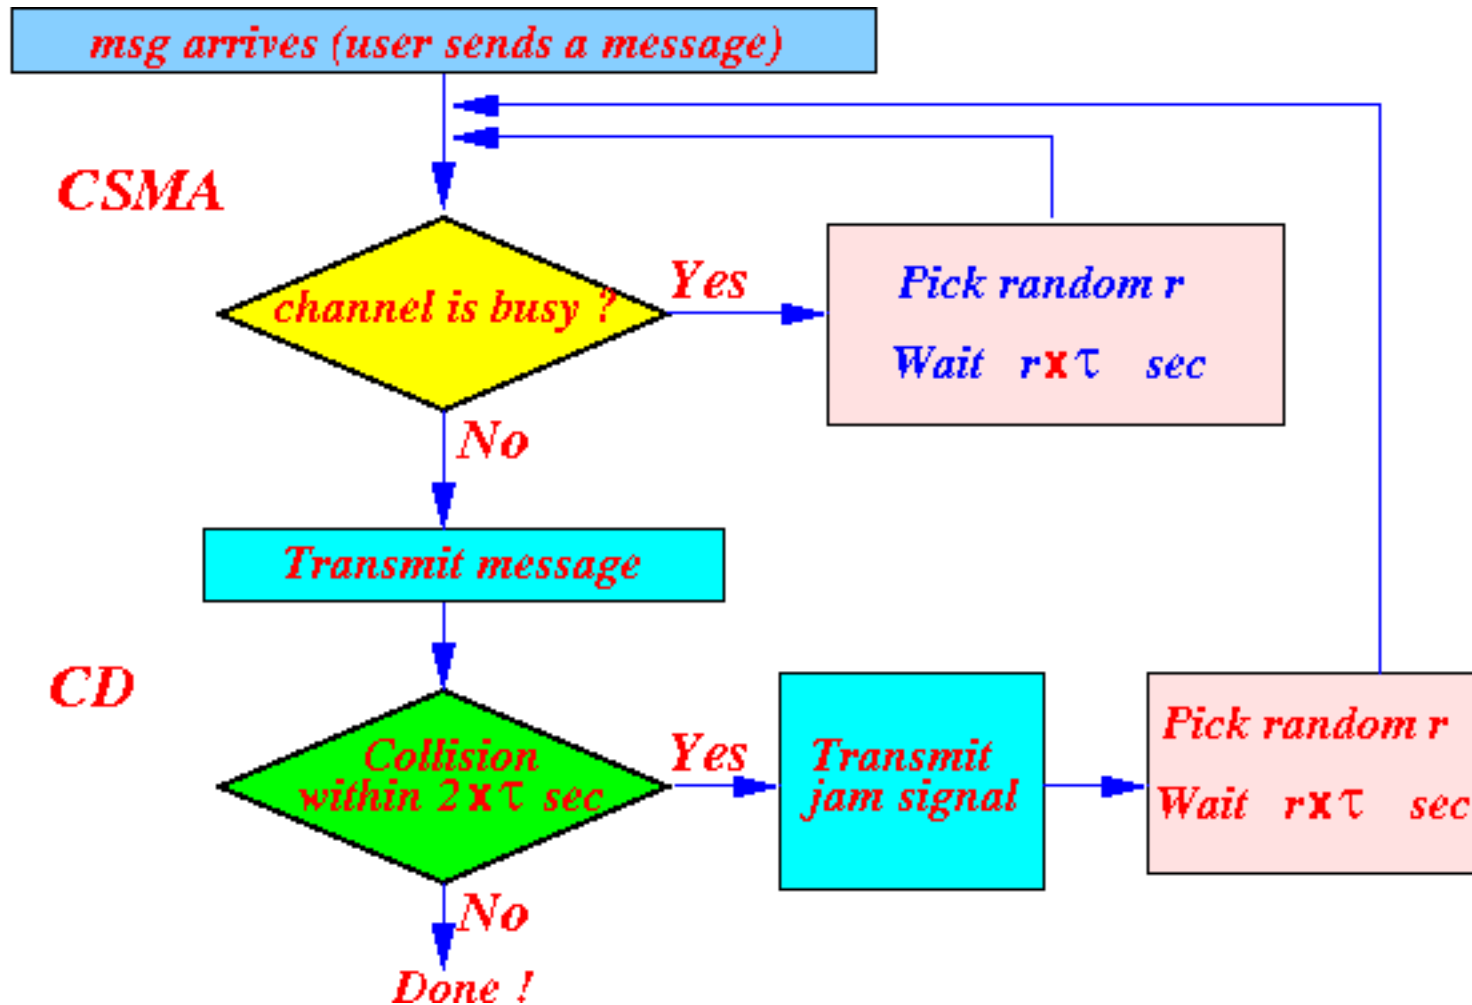
When we use the *non-persistent* CSMA backoff algorithm in the CSMA/CD flow chart:

-



we obtain the *Non-persistent* CSMA/CD protocol:

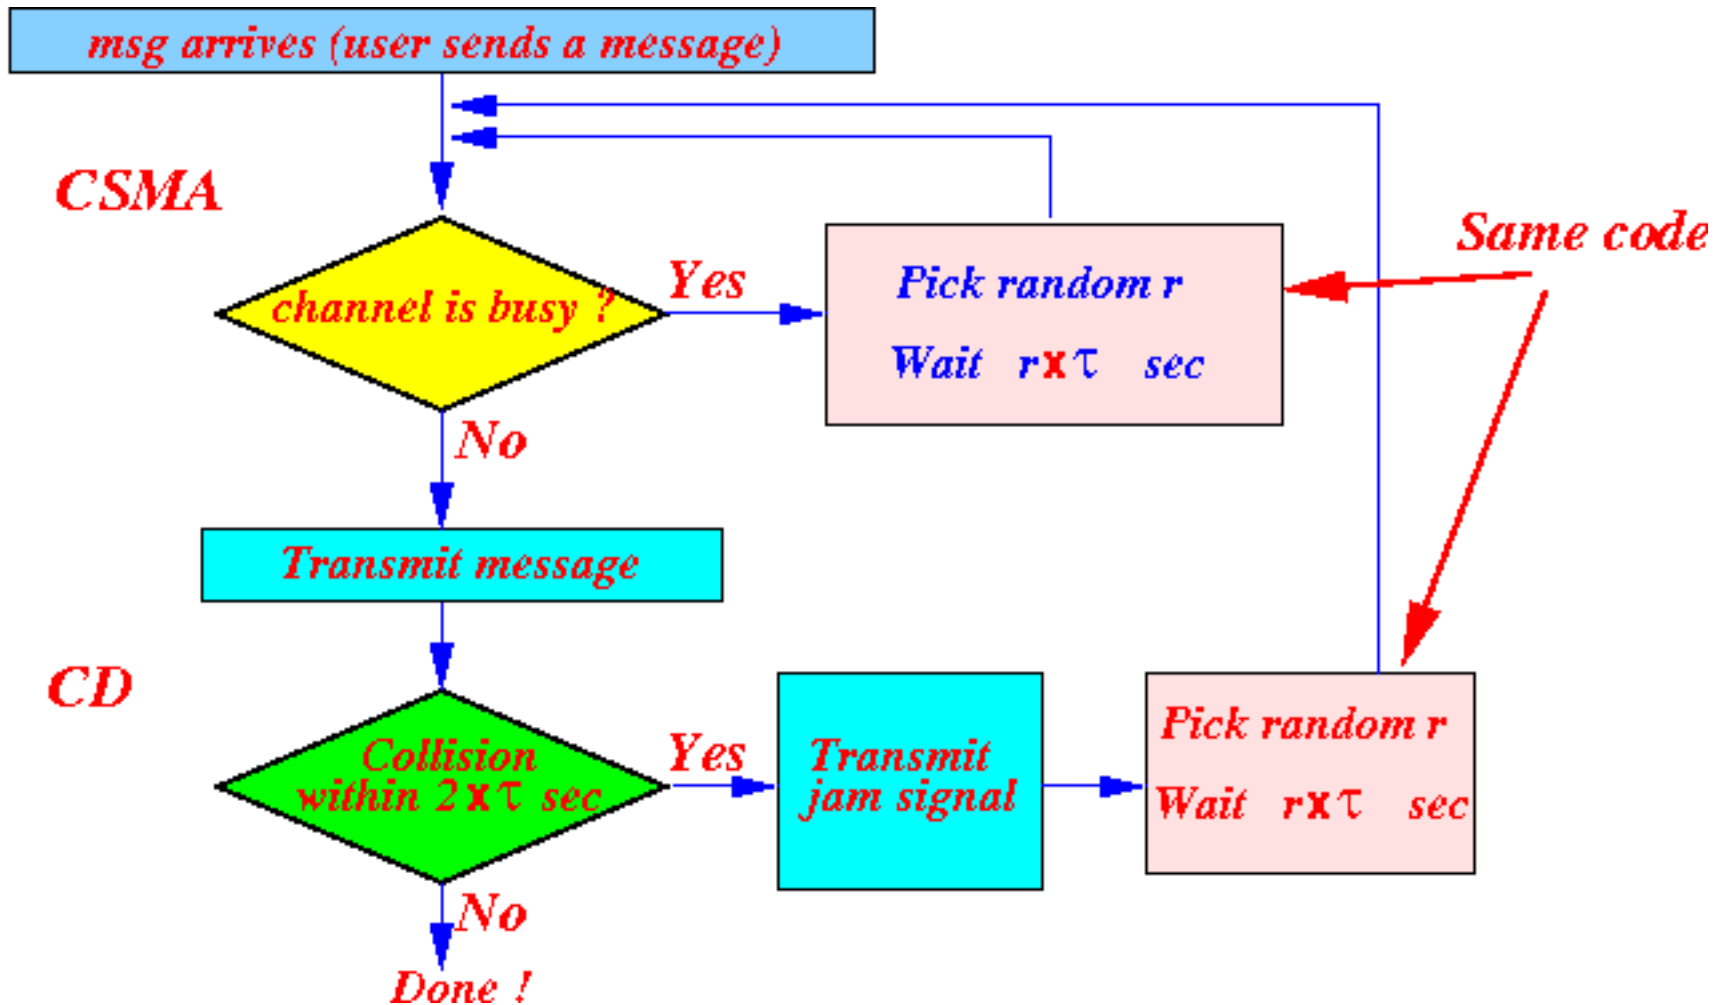




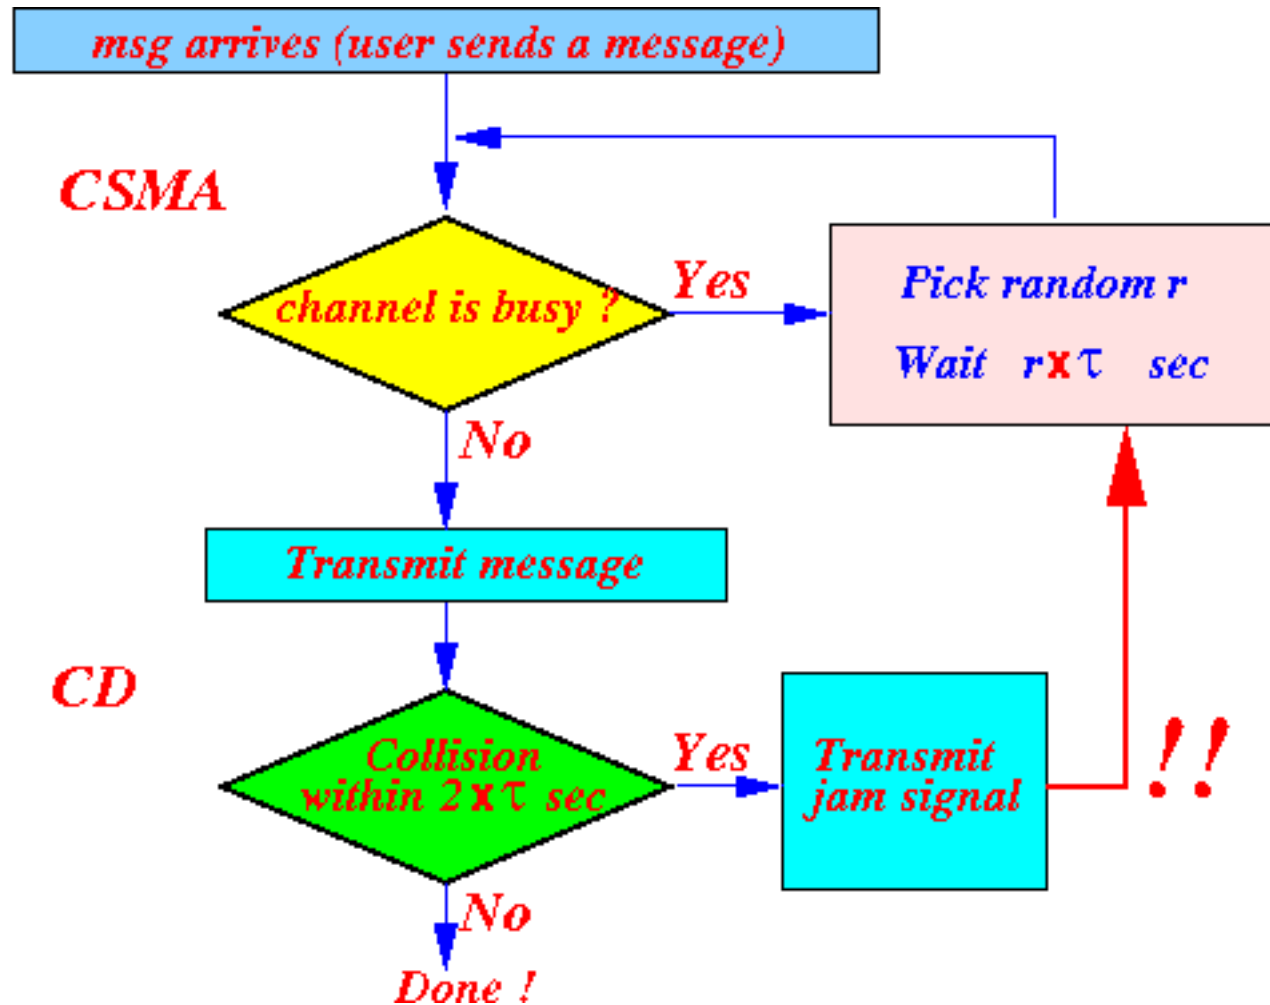
• τ = the end-to-end-delay



We can simplify the *previous* flow chart:

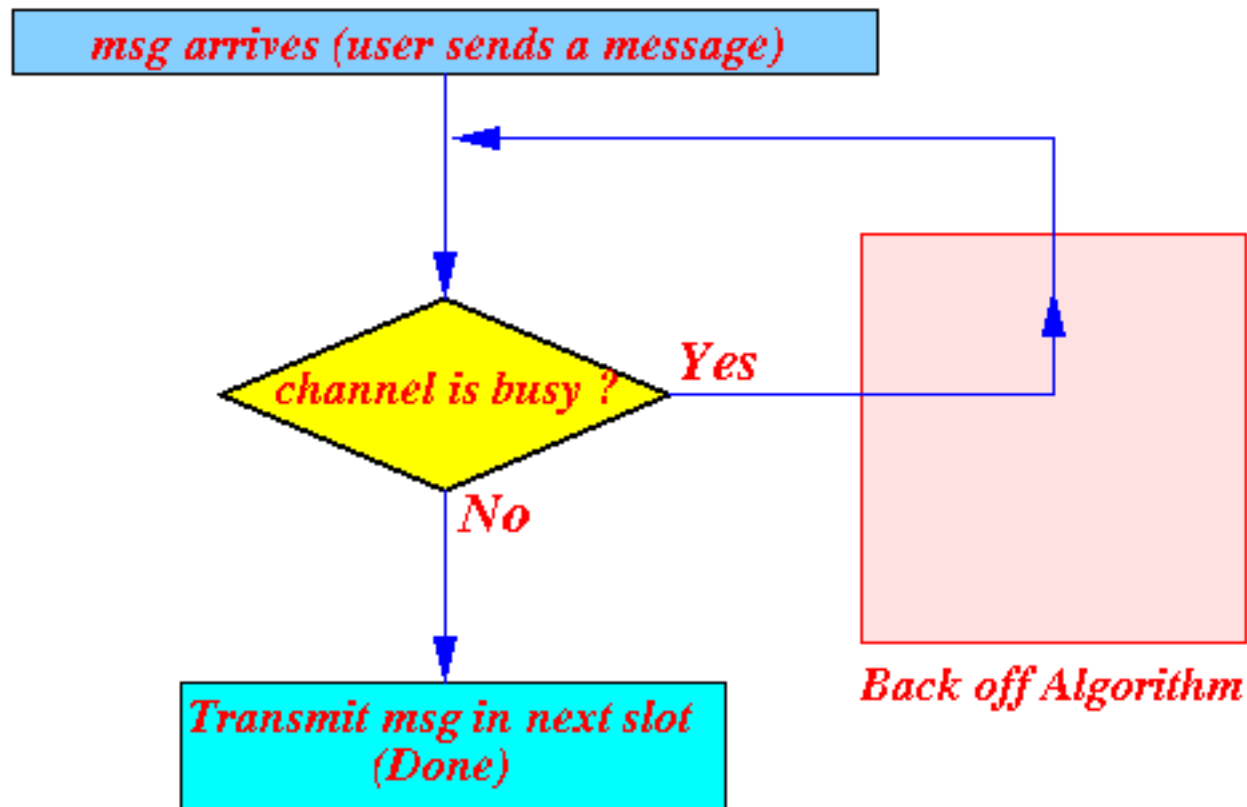


Simplify to the following

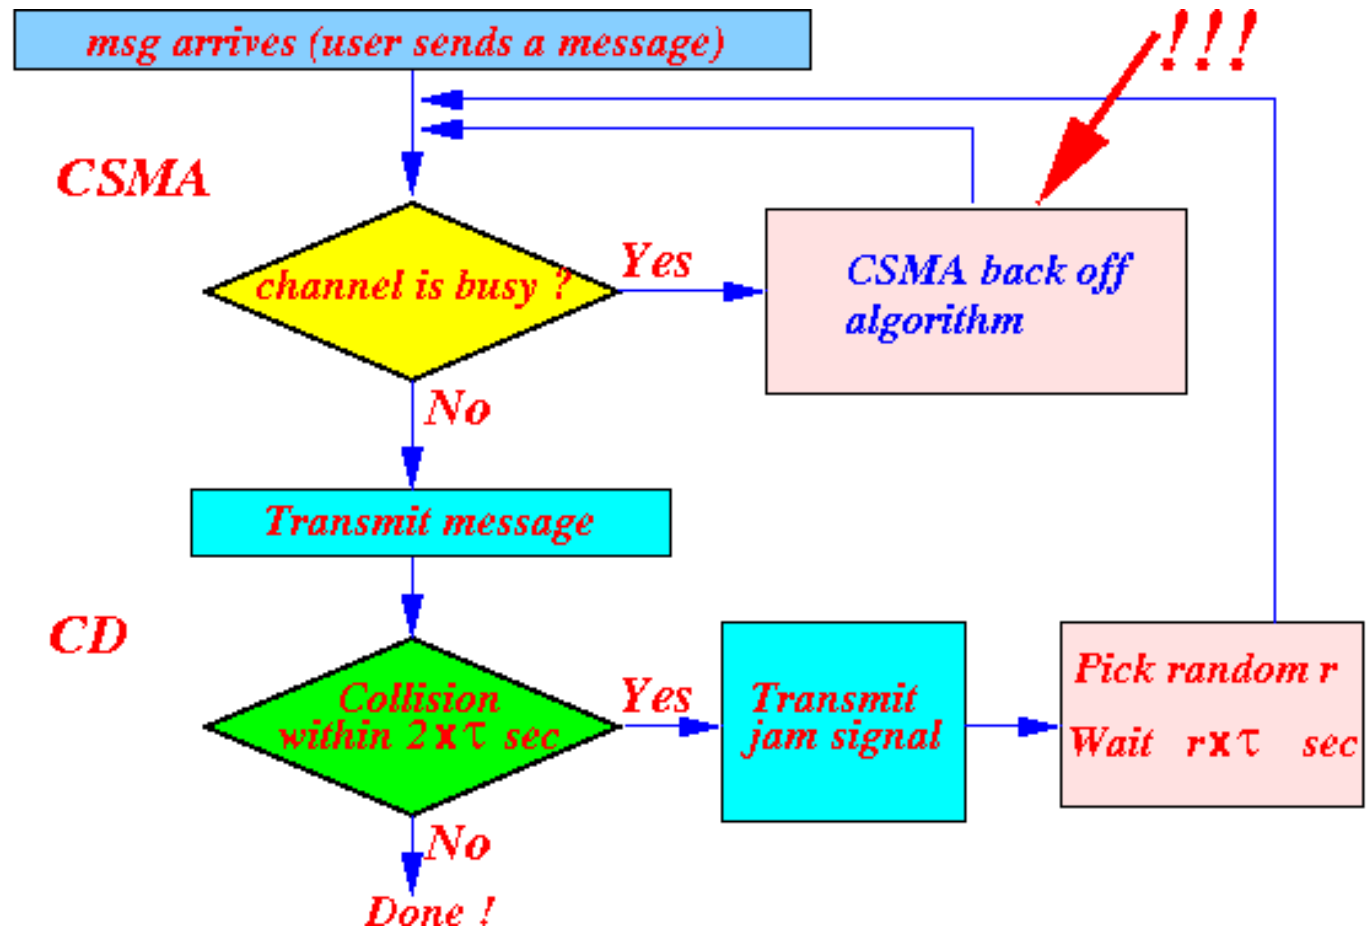


1-persistent CSMA/CD:

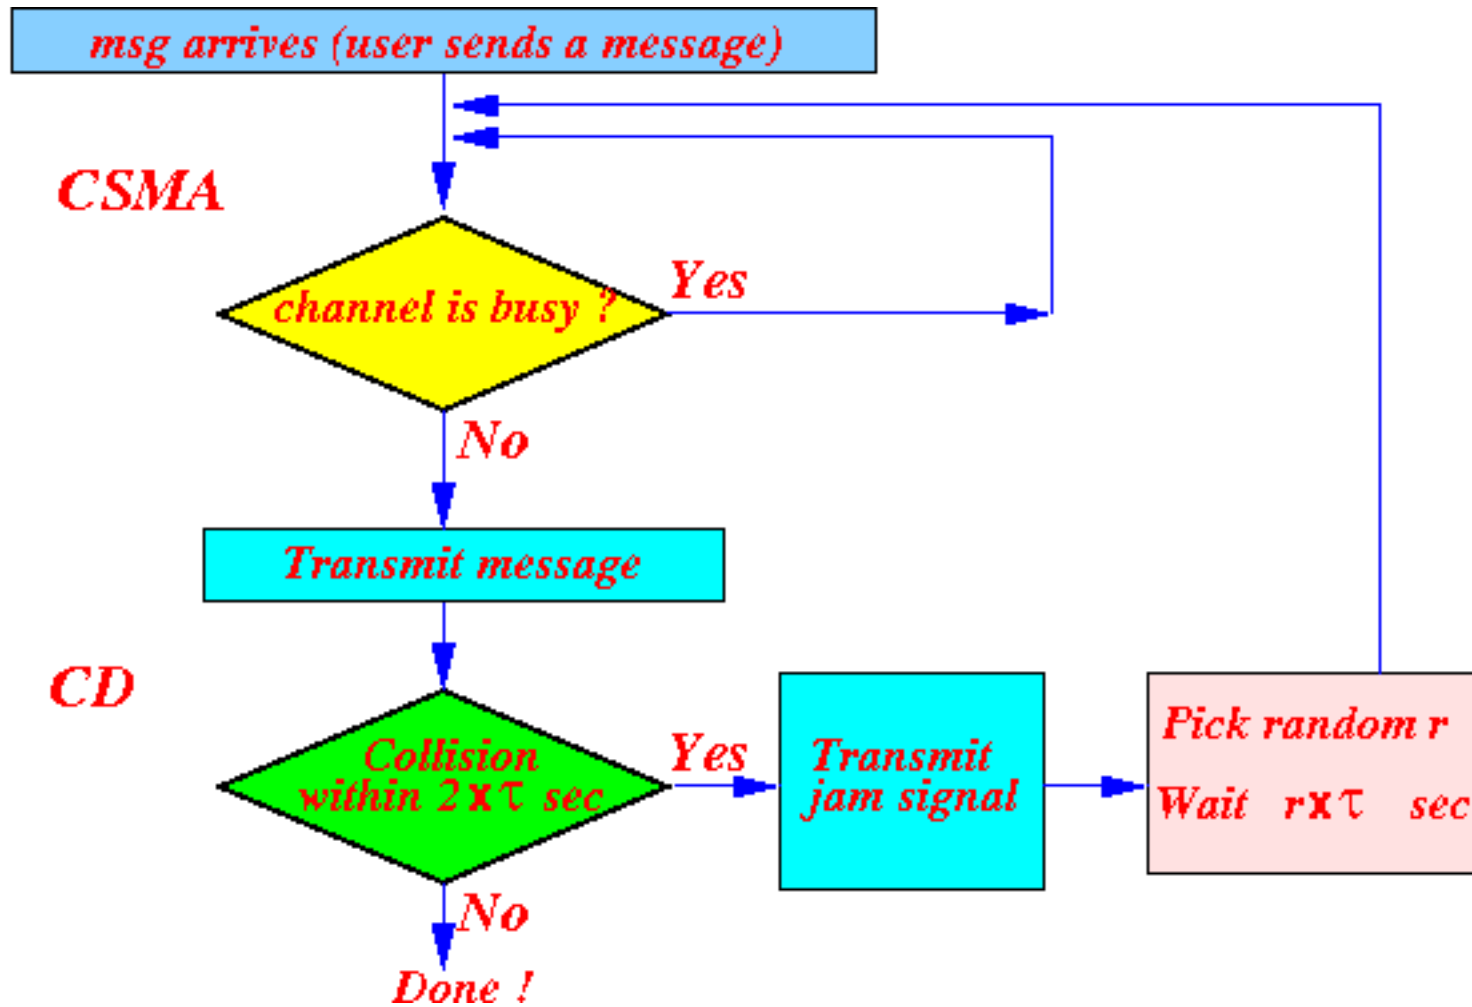
- Recall the *1-persistent* CSMA protocol:
-



- When we use the *non-persistent CSMA backoff algorithm* in the CSMA/CD flow chart:



The *1-persistent* CSMA/CD protocol:



Intro to Ethernet

- **Ethernet** was developed at **Xerox PARC** between **1973** and **1974**
- (PARC: **Palo Alto Research Centre** Incorporated is a branch of the **Xerox company**)
- The **Ethernet protocol** was invented by **Robert Metcalfe**
- **Metcalfe** was inspired by **Aloha** --- which he studied as part of his **PhD dissertation**



Documentation on the Ethernet

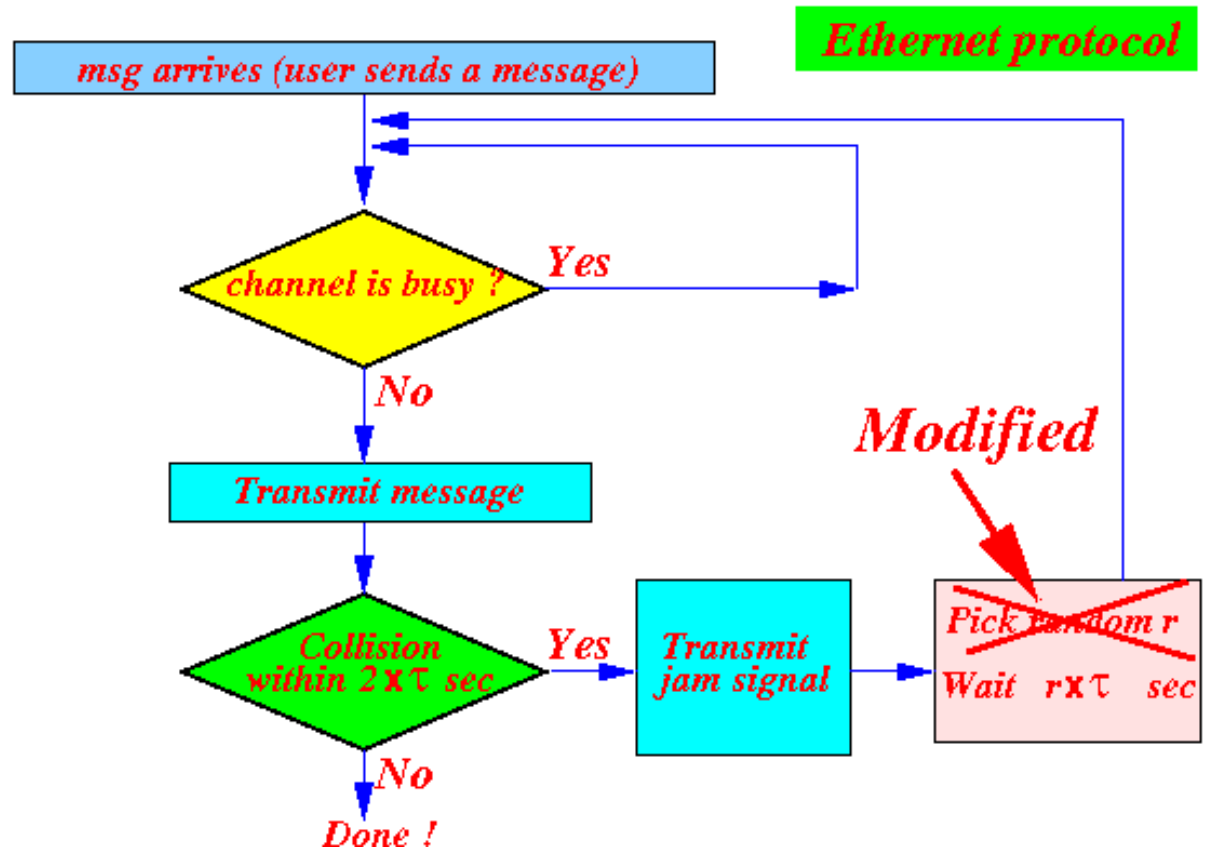
- "The Ethernet, A Local Area Network. Data Link Layer and Physical Layer Specifications" - **1980**
- **Standardization dates:**
- **International standard in 1983**
- **The IEEE 802.3 standard on: June 23, 1983**
-
-



The Ethernet *Protocol*

- A *modified 1-persistent CSMA/CD* protocol !!
- The Ethernet protocol flow chart:

Ethernet has a *modified* back off algorithm



The Ethernet *Frame* format



Meaning of the fields:

Preamble: discussed next

Start: 10101011 (Fixed pattern)

Src: Ethernet address of the sender

Dest: Ethernet address of the intended receiver

Length: Length of the data field

Data: Data field

CRC: Check sum



The *pre-amble*


- The **pre-amble** is used to *pad* the **Ethernet frame** so that the frame is **long enough** for **collision detection** purposes
- **Usage of the pre-amble:**
- A **transmitting node** can *only* detect a **collision** *while* the node is (still) **transmitting**
- The **pre-amble** will make **sure** that *simultaneous* senders will transmit **long enough** that the *simultaneous* senders will **detect a collision**



Example: how nodes can *miss* a collision


- Nodes *B* and *C* transmit *very* short frames:
-

B and C send a very short Ethernet frame



Their **transmissions** will **collide**:


Transmissions collide:



**However: the sending
nodes (*B* and *C*)
did not detect the collision:**




- **However:** the sending nodes (*B* and *C*) did **not** detect the collision:



Nodes *B* and *C* will *not* retransmit !

- The **pre-amble** will make sure that **transmitting node** will be transmitting *long enough* to detect a collision:
- **Nodes *B* and *C*** transmit *longer* frames: and hears the collision
- **B and C send longer Ethernet frame**



Ethernet Address

- *Each* nodes on a **Ethernet** is *identified* by a *unique* 48 bits **Ethernet** address
- The **Ethernet** address is *permanently* inscribed in the **Ethernet** card of a computer
-



- **UNIX command** to find the **Ethernet Address** on a **computer**:
- \$ ifconfig -a

- eth0
- Link encap:Ethernet HWaddr 2C:41:38:8B:05:38
- inet addr:170.140.150.36
- Bcast:170.140.151.255 Mask:255.255.254.0
- inet6 addr: fe80::2e41:38ff:fe8b:538/64
- Scope:Link UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
- RX packets:681868507 errors:0 dropped:0 overruns:0 frame:0
- TX packets:628584348 errors:0 dropped:0 overruns:0 carrier:0 collisions:0 txqueuelen:1000
- RX bytes:335075333658 (312.0 GiB) TX bytes:540484678811 (503.3 GiB)
- Interrupt:20 Memory:fb000000-fb020000 lo Link encap:Local Loopback inet addr:127.0.0.1
Mask:255.0.0.0 inet6 addr: ::1/128 Scope:Host UP LOOPBACK RUNNING MTU:16436 Metric:1
RX packets:267964 errors:0 dropped:0 overruns:0 frame:0 TX packets:267964 errors:0
dropped:0 overruns:0 carrier:0 collisions:0 txqueuelen:0 RX bytes:107815976 (102.8 MiB) TX
bytes:107815976 (102.8 MiB)



Explanation:

- Ethernet address = 2C:41:38:8B:05:38
(Hexadecimal digits)
- 2 C 4 1 3 8 8 B 0 5 3 8 =
- 0010 1100 0100 0001 0011 1000 1000 1011 0000
0101 0011 1000



Why are there *no* send/receiver sequence numbers ???



There are:

- *No Send* sequence number
- *No ACK* sequence number

in Ethernet frames



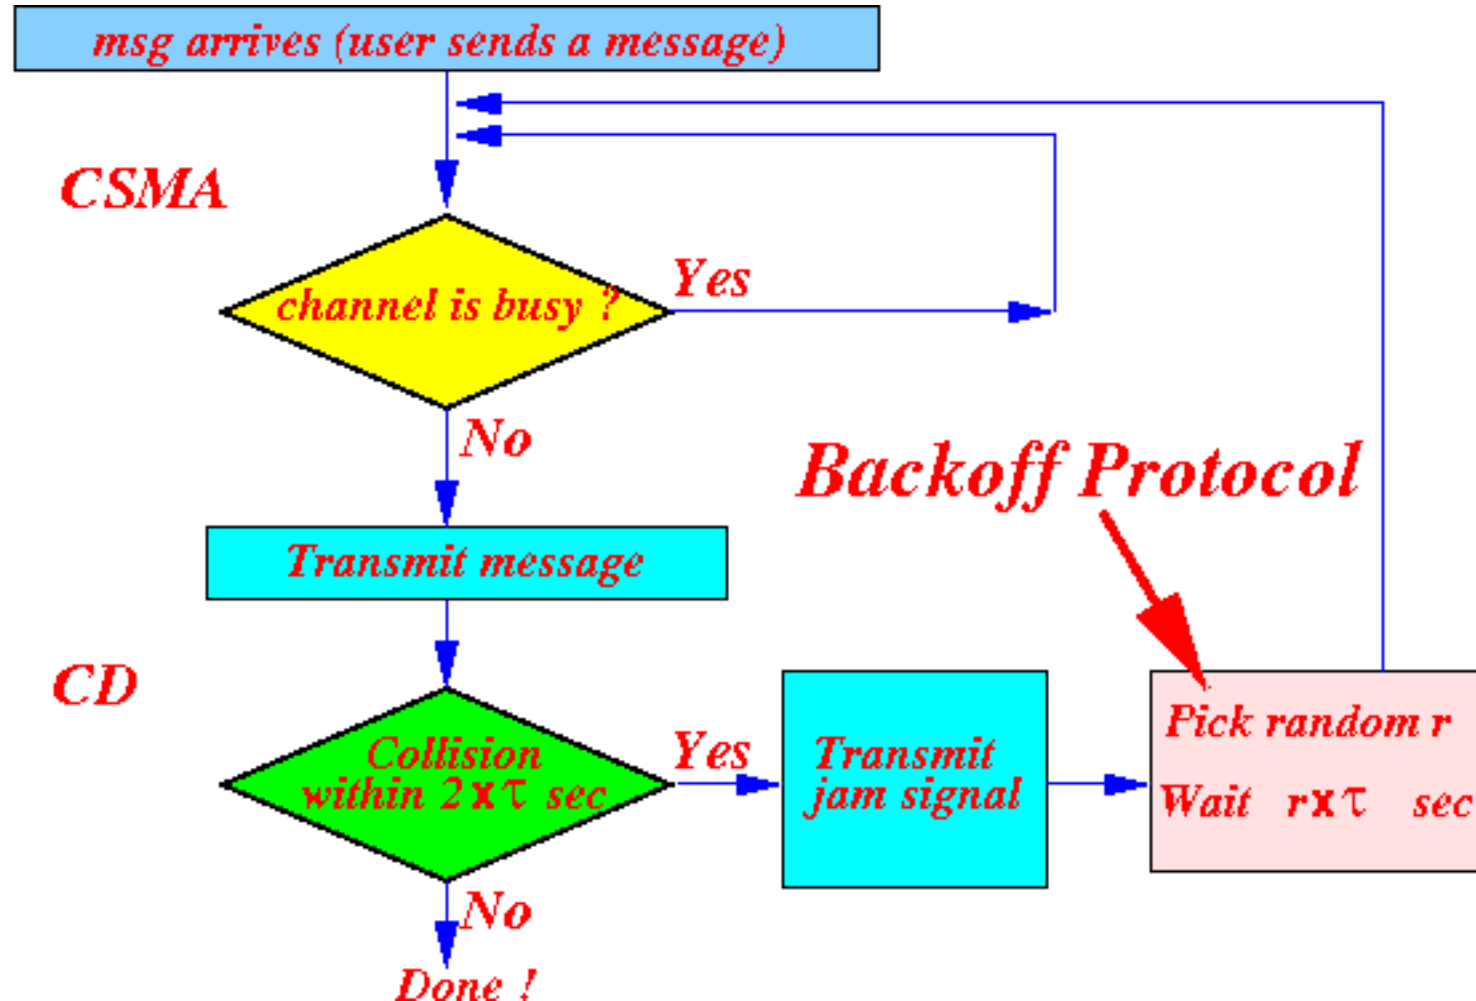
Reason:

- The **Ethernet (LAN) protocol** does *not* provide **reliability** function !!
- **Ethernet** relies on the *Transport Layer* to ensure **reliable communication**
- There is an **ARQ protocol** in the *Transport layer* !



The *Exponential* Backoff Algorithm of the Ethernet

- The Backoff Protocol

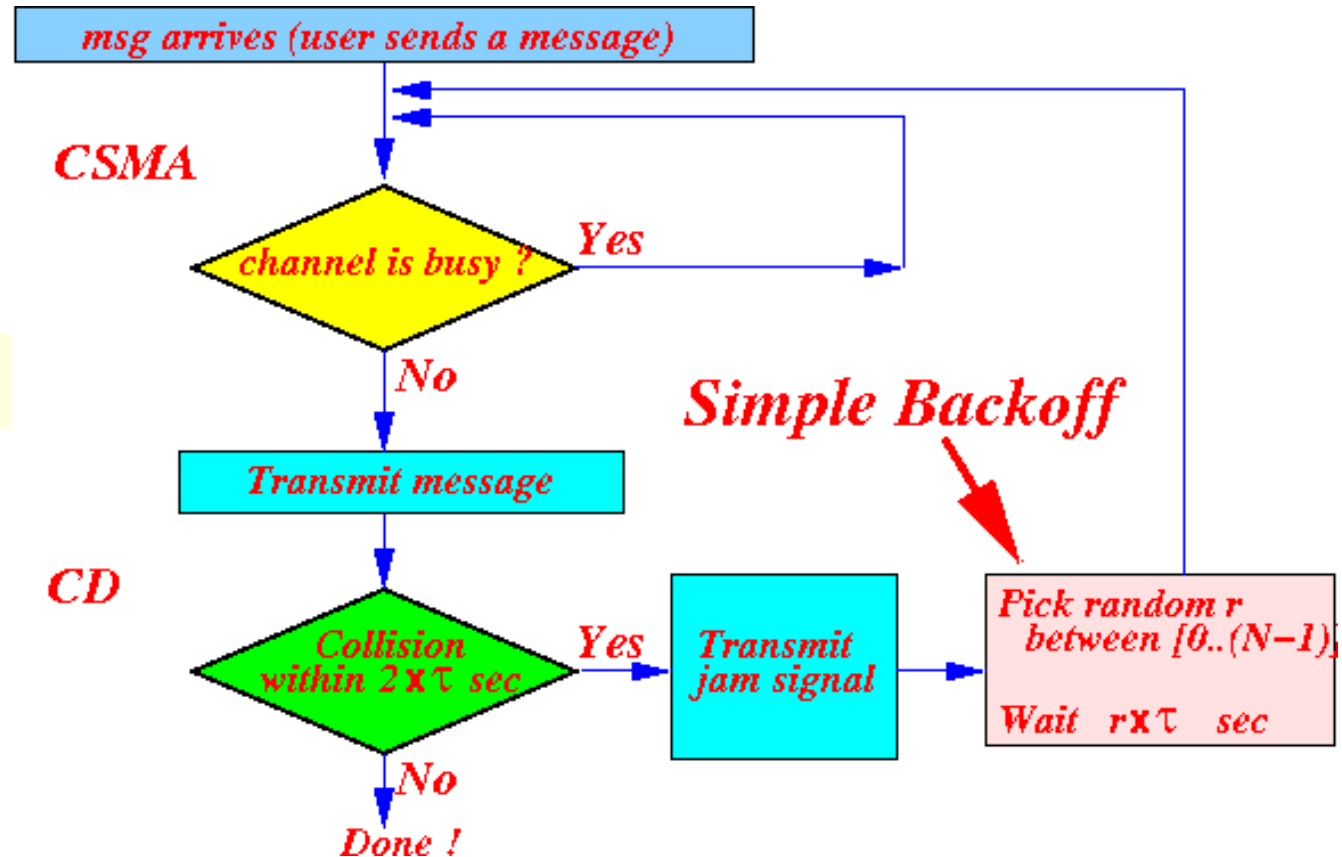


- The **backoff protocol** is **executed** when a **node have (already) detected a collision !!!**
- *Multiple* nodes are *ready* to **transmit !!**
- (*One* **transmission** can *never* result in a **collision.....**)
- *Goal* of the **Backoff** protocol:
- *Re-schedule* the **transmissions** of the *collided* nodes so that:
- the **likelihood (chance)** of *subsequent* collisions is **minimized**



A simple-minded backoff protocol

- Pick a random number r between $[0..(N-1)]$ (N is fixed)



- Wait $r \times \tau$ time




Weakness of the *simple-minded* backoff algorithm

- If N is *small*, and the system is very busy (= *many* nodes have *collided*),
- then it will be *very* likely that *more* than one node will select the *same* random number


They all wait for current transmission to end....

Result : frequent collisions



- On the other hand, if N is *large*, but the system is *lightly* loaded, then:
- Nodes *may* select an *large* random value
-

They all wait for current transmission to end....



-
- **Result:**
 - The **node** will **back off** for a *long* time (for nothing)...
 - **Result:**
 - **Wasting of bandwidth...**
 - **Conclusion**
 - We need a *dynamic* backoff algorithm that can *adjust* to *different* network operating conditions !!!

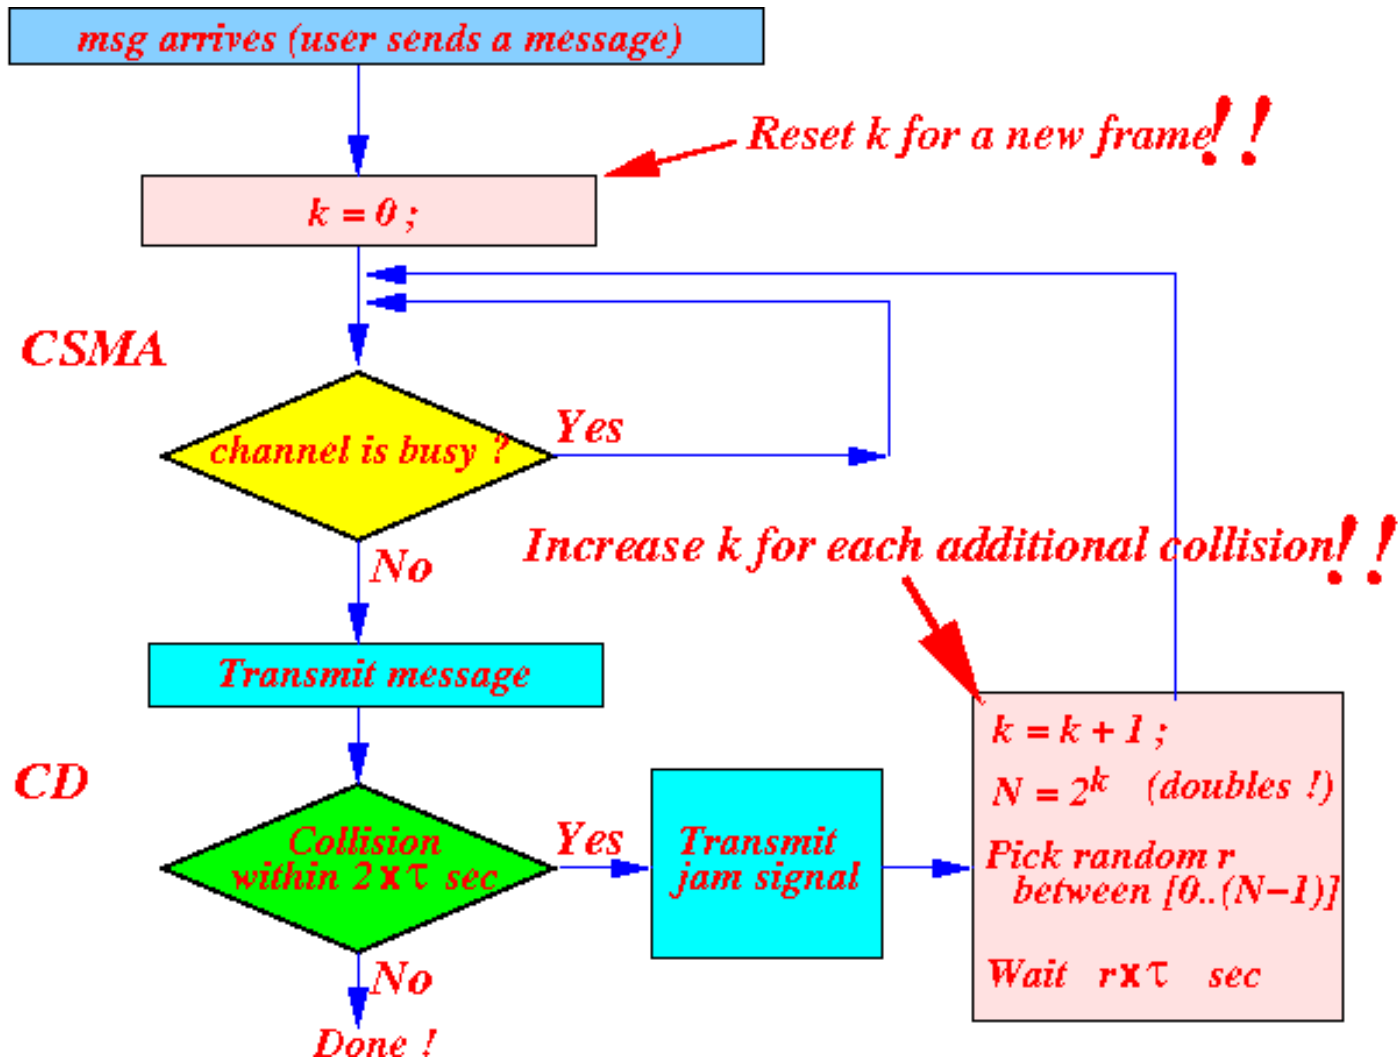


The Exponential Backoff Protocol

- *Adaptive* back off:
- Allow **collided nodes** to pick *small* back off periods in *earlier* re-transmission attempts
- And pick *larger* back off periods when there are *more* re-transmission attempts
- The exponential backoff algorithm flow chart:
-



The exponential backoff algorithm flow chart:



Explanation

- When a **frame** is **transmitted** for the *first* time:
- Reset $k = 0$
- *If* there was a **collision**, the *first* retransmission attempt will use:
 - $k = k + 1 = 0 + 1 = 1$
 - $N = 2^k = 2^1 = 2$
- For the *first* re-transmission attempt, the **node** will pick a **random number** r from the range:
 - $r \in [0, 1]$
- This **range** is **suitable** when there are *very few* (about 2) **nodes** involved in the **collision**



- It the **node** is **involved** in *another* **collision** while trying to **transmit** the *same* **frame**, then:
- $k = k + 1 = 1 + 1 = 2$
- $N = 2^k = 2^2 = 4$
- So the *second* **retransmission attempt** for the *same* **frame** will use a **random number** r from the range
- $r \in [0..3]$
- The **node** is *adjusting* to **accommodate** for a *larger* **number of nodes** (i.e., a *heavier loaded* **situation**)



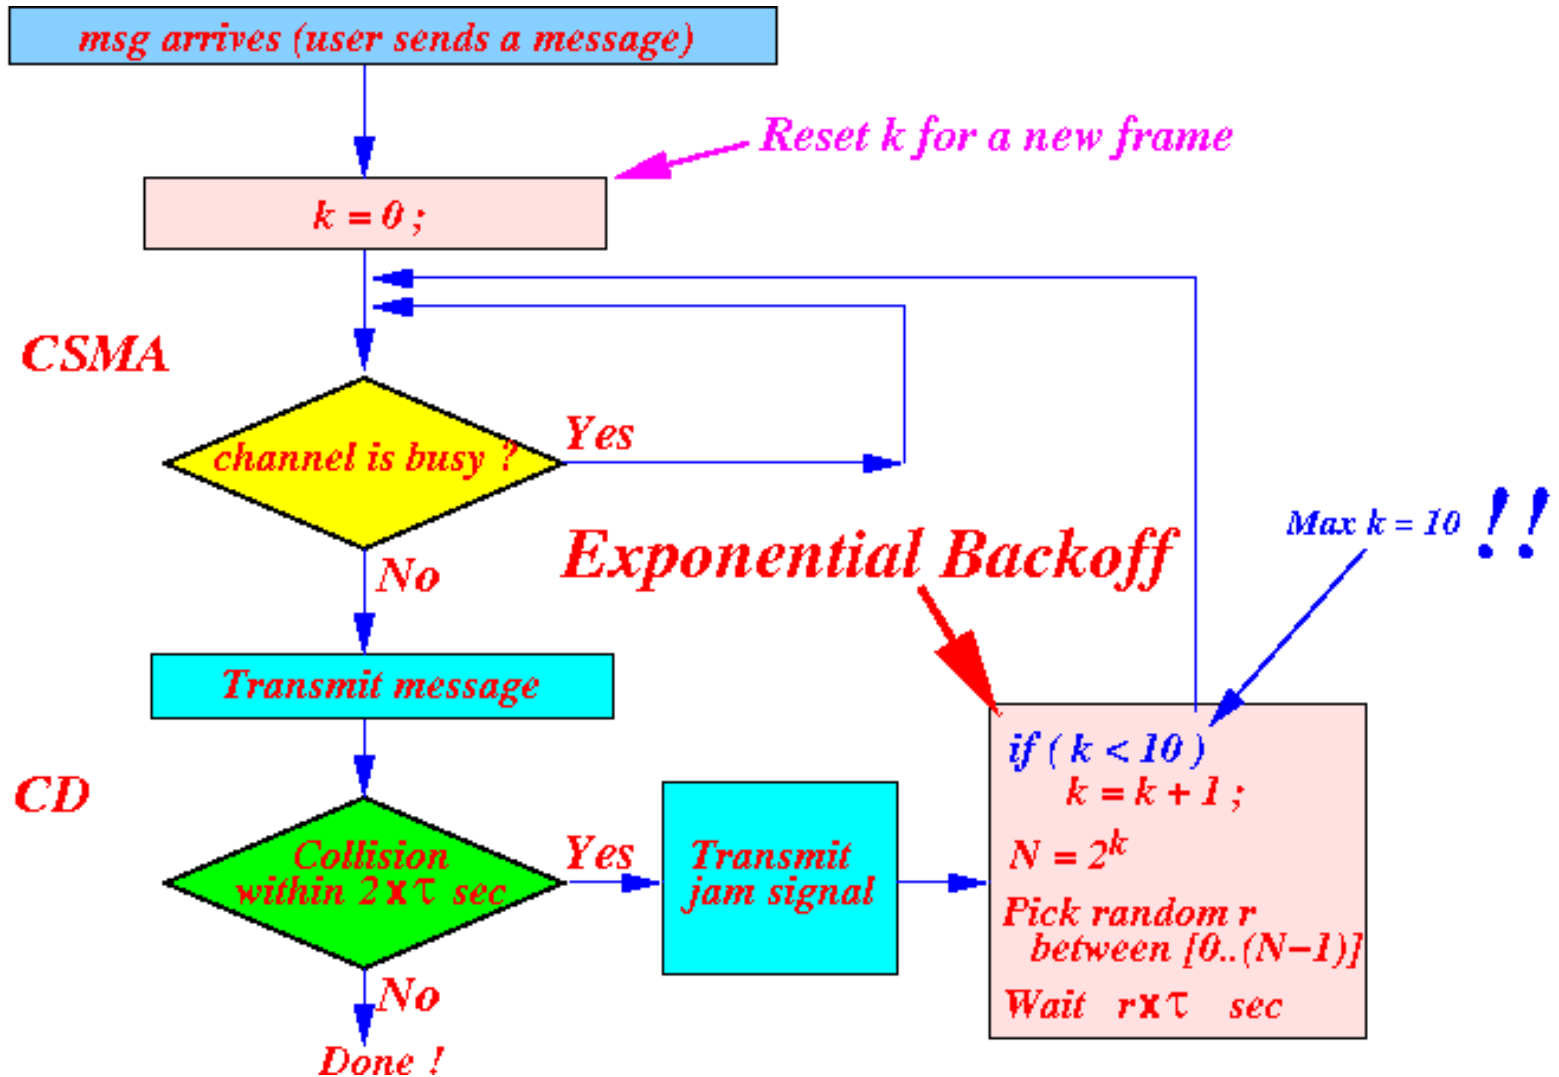
-
- If there is **another collision**, the *third* retransmission attempt of the *same* frame will use:
 - $r \in [0..7]$
 - And so on...



-
- A *small* adjustment: upper bound on random range
 - One Ethernet network can support at most 500 node
 - Therefore
 - The maximum range from which the random numbers are selected is:
 - $[0..2^{10}-1] = [0..1023]$



The *final* version of the Ethernet Medium Access Control (MAC) algorithm:




The *unfairness* of the Ethernet Backoff Algorithm

- Channel capture effect
- Channel capture effect = a phenomenon where one user of a *shared* (= broadcast) medium "*captures*" (= hogs) the medium for a significant time
- *Unfairness* of the Ethernet's backoff algorithm
- Suppose:
 - The nodes *A* and *C* in an Ethernet has a *large* number of frames to transmit
 -




- **Consider the following scenario:**
- **Nodes *A* and *C* transmits at the *same* time:**

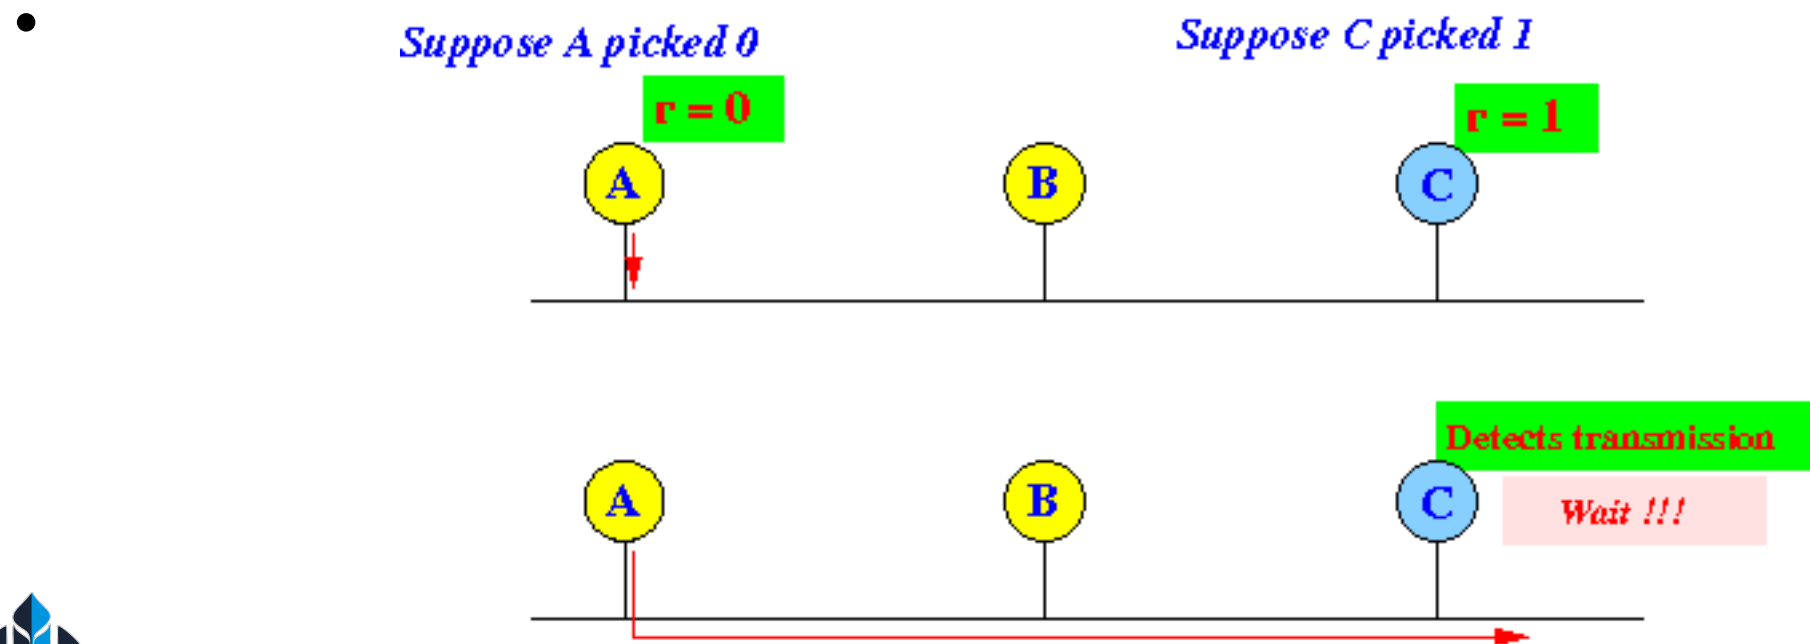


- The **transmissions** of **nodes *A* and *C*** will **collide** and they will **back off**
-

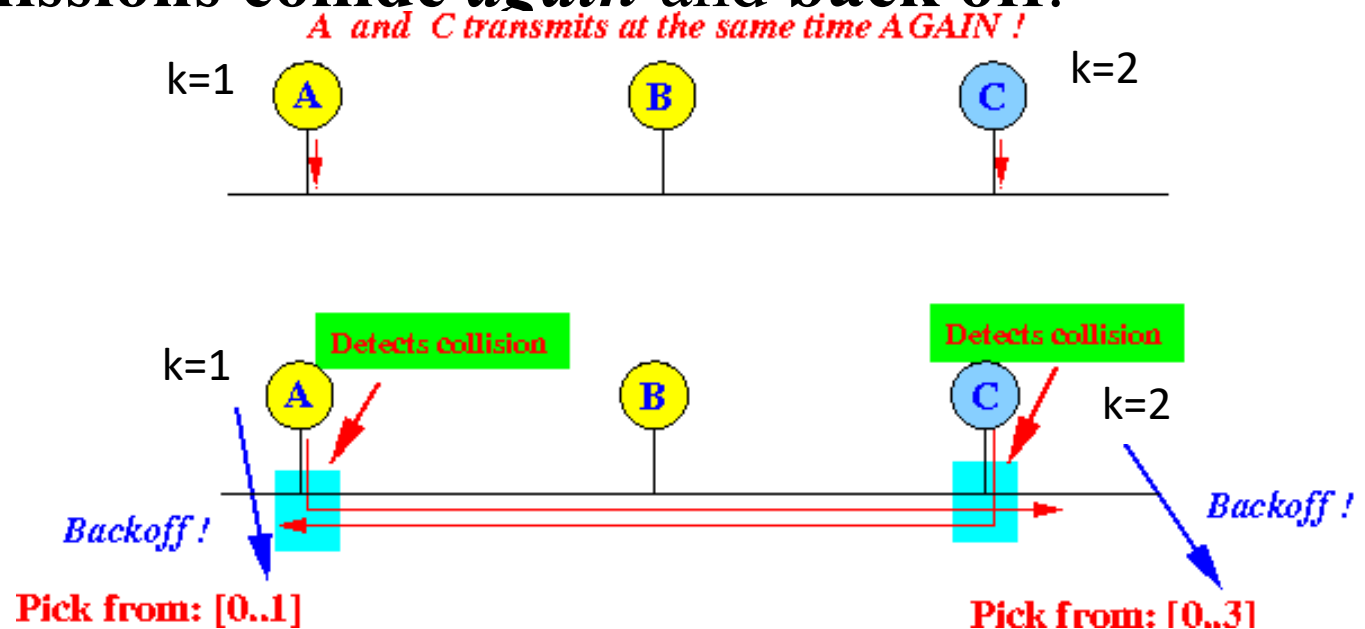
- The **transmissions** of nodes **A** and **C** will **collide** and they will **back off**



- **Suppose:**
- Node *A* picked 0
- Node *C* picked 1
- **Then node *A* will transmits *first*:**



- Node *A* will **complete** its **transmission**
- When node *A* is **done**, *A* will **transmit *again***
- *Then* node *C* will *also* **transmit**.
- Their **transmissions collide *again*** and **back off**:



Notice that:

- **Node A** transmits the **frame** for the *first* time, so:
- Node A will pick a **random number** from the range **range [0..1]**
- **Node C** transmits the **frame** for the *second* time, so
- Node C will pick a **random number** from the range **range [0..3]**
- **Problem**
- The *likelihood* that **node A** will transmit *before* the node **C** is *higher* !!
- The **probability** is **calculated** in the **section below..**
- This is *unfair* for **node C** !!!



Unfairness calculation....

- **Question:**
- If **node A** picks from **[0..1]** and **node C** picks from **[0..3]**, then:
- What is the **probability** that node **A** will transmit *before* node **C** ?
- **Pre-question:** *when* will node **A** transmit **before** node **C**:
- Node **A** transmits **before** node **C** if
- Node **A** picks a *smaller* (**random**) **value** than node **C**



Answer

- A will transmit *before* C when
- A picks 0 and C picks 1
- A picks 0 and C picks 2
- A picks 0 and C picks 3
- A picks 1 and C picks 2
- A picks 1 and C picks 3
- The probability that C will pick any one number is 0.25 The probability that B will pick any one number 0.5
- The chance that C will transmit before B is:
- $$\frac{(0.25 \times 0.5) + (0.25 \times 0.5) + (0.25 \times 0.5) + (0.25 \times 0.5) + (0.25 \times 0.5)}{0.5} = 0.625$$



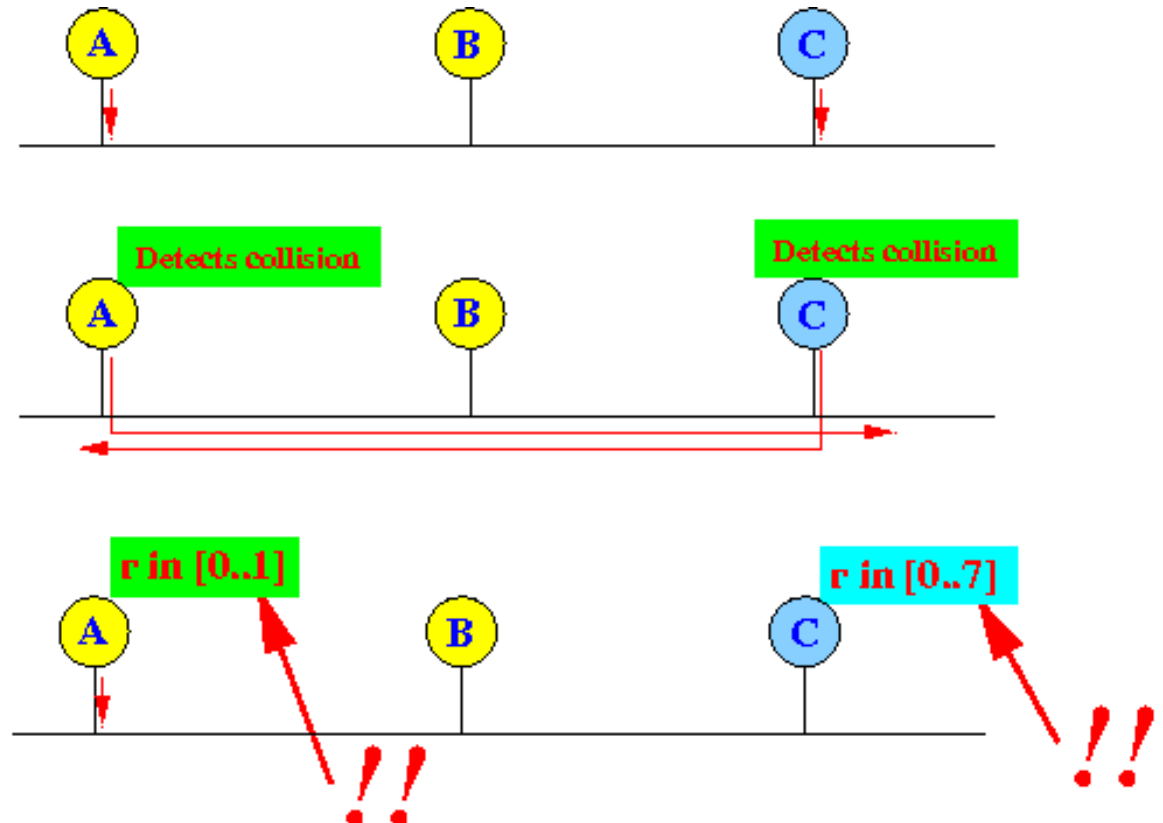
Summary:

- If **node A** picks from **[0..1]** and **node C** picks from **[0..3]**, then
- The **probability** that node **A** transmits *before* node **C** = **0.625**
- The **probability** that node **C** will transmit *before* node **A** = **0.125**
- **Conclusion**
- This is *not* fair !!



And the **bad news** is:

- It will get **even worse** if **C** loses **again** ...
- **Suppose** node **B** transmits first and during **B's** **transmission** the node **A**, scheduled a transmission.



-
- The **likelihood** that **node C** will **transmit first** is now **very very *small* !!!**
 - **NOTE**
 - This **problem** is **very well-known**, and it is called the
 - **Ethernet Capture Effect**

